

A Deep Dive into Model Structure: Do Deep Learning Models Learn Verifiable and Transferable Hydrological Knowledge?

Jiachen Kong¹, Zhaocai Wang^{1,*}, Zhaoyang Zhu¹

¹Shanghai Ocean University, Shanghai, China

*Corresponding author: zcwang1028@163.com

Abstract

With the widespread use of deep learning for streamflow prediction, a central question remains unresolved: do these models learn hydrologically meaningful knowledge, or only predictive correlations? Existing studies mainly rely on post-hoc methods such as SHAP and LIME, which provide descriptive attribution but offer limited evidence on verifiability, behavior-level physical consistency, and cross-model transferability. To address this gap, we propose Sparse Transformer-based Explainable Hydrology (STEH), a framework that treats interpretability as a primary objective. STEH links inputs, internal circuit structures, and outputs through structured sparsification and pathway-level analysis, and then evaluates learned knowledge with circuit ablation, sensitivity experiments, and transfer tests. Across large-sample basins, results show that structurally simpler models tend to form clearer decision pathways, whereas excessive complexity introduces redundancy and can reduce predictive performance. We further find, in probe-based analyses, that physically meaningful signals are broadly present across models with different accuracy levels, and that STEH-derived structural knowledge can improve a Random Forest benchmark after transfer. These findings suggest that deep learning can support both prediction and hydrological knowledge discovery, and motivate a shift from descriptive explanation to systematic verification in hydrological interpretability research.

Keywords: Model interpretability, Explainable artificial intelligence, Structural sparsity, Hydrological modeling, Streamflow prediction, Transformer

1 Introduction

Streamflow prediction plays a central role in water resource management, flood mitigation, and watershed operation regulation [1]. In the context of increasing hydrological and climatic uncertainties [2], it is a key support for high-stakes decision-making in areas like optimal water allocation and precise flood risk management [3]. Therefore, the reliability and accuracy of the model are of equal research importance. In recent years,

machine learning and deep learning models have demonstrated excellent performance in streamflow prediction [4, 5], but the inherent "black-box" nature of these models makes their internal decision-making mechanisms difficult to interpret [6], significantly limiting the transparency and reliability of the predictions [7], and potentially misguiding actual decision-making. Thus, improving model interpretability while maintaining predictive performance has become a core challenge in hydrological modeling [8], garnering widespread attention and research from the academic community.

To address the "black-box" characteristic of data-driven models, some scholars have attempted to integrate physical knowledge of hydrological processes into neural network architectures to enhance the physical consistency of the models [9]. Among these, Physics-Informed Neural Networks (PINN) are one of the most representative approaches [10], with the core idea of embedding hydrological control equations into the model's loss function as constraints, so that the model strictly adheres to physical laws while fitting observed data [11, 12], thereby improving the physical plausibility of the model outputs. However, physical consistency does not fundamentally resolve the issue of model interpretability [13]. The model's utilization mechanisms for input variables [14], as well as the core logic it uses to represent hydrological processes internally, still remain largely ambiguous [15], making it difficult to meet the deeper interpretability requirements for practical decision-making.

Another mainstream approach relies on post hoc analysis techniques to explore model interpretability. Among them, Local Interpretable Model-agnostic Explanations (LIME) [16] and SHapley Additive Explanations (SHAP) [17] are the two most widely used methods. These techniques build local surrogate models to approximate predictions or decompose them based on cooperative game theory [18], quantifying the contribution of each input feature to the model's output. In this way, they partially break the "black-box" barrier of model transparency [19], providing intuitive tools for analyzing the impact mechanisms of key hydrological variables such as precipitation and evapotranspiration. However, these post hoc explanation methods have inherent limitations: the explanation results heavily depend on local approximations or statistical decomposition processes, which are easily influenced by data distribution characteristics and parameter settings. Furthermore, these methods lack deep integration with hydrological processes, making it difficult to fully reflect the intrinsic logic behind the model's decision-making, and thus the reliability of the explanation results remains to be improved.

Overall, both PINN methods and post hoc explanation techniques have made significant progress in enhancing the interpretability of streamflow prediction models, offering feasible approaches to address the "black-box" problem. However, both methods still have some shortcomings that can be further improved, which lead to concerns when applying interpretability results to guide high-risk watershed management decisions. Specifically, these concerns manifest in three main areas:

First, the interpretability results lack clear verifiability [20]. Unlike model predictions, which can be directly compared and validated against observed streamflow data, feature attribution-based interpretability results do not have objective ground truth values, making it difficult to quantitatively verify them through independent experiments [21]. Existing studies have shown that even if a model's prediction accuracy is high, its interpretability results may only reflect complex statistical fitting relationships or spurious correlations [22]. Relying solely on feature attribution analysis makes it challenging

to determine whether the model has genuinely learned physically meaningful hydrological relationships [23], which severely limits the trustworthiness of interpretability results in high-risk decision-making environments.

Second, there is a lack of systematic evaluation of the model’s behavior. Most current interpretability studies remain confined to feature-level attribution analysis based on SHAP or LIME [24, 25], without thoroughly testing whether the model’s response to changes in hydrological conditions aligns with physical logic [26]. In practice, a physically reliable streamflow prediction model should generate consistent physical responses when key driving factors are disturbed [27]—for example, streamflow should steadily increase when precipitation significantly rises [28]. Without such perturbation testing or counterfactual validation [29, 30], interpretability results cannot be effectively linked to actual hydrological behavior [31], and thus fail to provide genuinely valuable insights for decision-making.

Third, there is a lack of systematic evaluation of the transferability of interpretability results. Current interpretability analyses are mostly confined to a single model architecture and rarely focus on whether the extracted hydrological relationships remain valid across different model families [32]. If identified feature relationships cannot be reproduced or applied in other models, these relationships are more likely to reflect the parameterization characteristics of a specific model, rather than generalizable hydrological scientific knowledge. Therefore, cross-model validation is crucial for establishing the scientific value of interpretability results [33]. It is worth noting that in other fields, systematic interpretability research paradigms have gradually advanced, and standardized paradigms can significantly enhance trust in model results. However, such attempts remain relatively limited in the hydrological modeling domain, and a mature interpretability verification system has yet to be formed.

This study systematically uncovers the internal information flow mechanisms of deep learning-based streamflow prediction models through multidimensional verification methods, confirming the presence of physical knowledge in the model and clarifying that it is the core foundation for achieving accurate streamflow predictions. Based on this, the study focuses on the following key scientific questions:

1. What is the intrinsic relationship between model structure and prediction performance?
2. Does physical knowledge exist in deep learning-based hydrological models, and what patterns govern its spatial and feature distributions?
3. Can the model generate feature representations that adhere to physical laws for different hydrological scenarios?
4. Does the physical knowledge learned by the model possess transferability across scenarios and model structures?

To address these scientific questions, this paper proposes an interpretable hydrological model (STEH), aiming to achieve full-chain interpretability from input features, internal model structures, to output results. The main contribution of this study lies not only in the development of an interpretable deep learning framework, but also the

establishment of a reference path for hydrological model interpretability verification. Framed as *A Deep Dive into Model Structure*, this work integrates structural analysis, behavioral verification, and cross-model evaluation to move interpretability beyond descriptive attribution toward rigorous assessment of whether model knowledge is *verifiable*, physically meaningful, and *transferable*. In doing so, it provides preliminary insights into a fundamental question: *what kind of hydrological knowledge is learned, and is it verifiable and transferable?*

2 Results and discussion

Next, we conduct an interpretability study of deep learning-based streamflow prediction models using the CAMELS-US dataset. After screening, 521 study sites are selected (details are provided in the Section 4.3), and the model performs well in the majority of regions (Supplementary Information 1). By integrating the model circuitry with various interpretability methods and perspectives, we provide an in-depth explanation of how the proposed interpretable deep learning framework (STEH) better captures the internal knowledge representation and physical consistency issues that existing models struggle to resolve.

2.1 Interpretable Circuit Structures Across Hydrological Stations

Given the large number of stations (521), it is impractical to visualize the learned circuits for all stations individually. Therefore, a representative sampling strategy is adopted to systematically examine circuit structures under different predictive conditions.

Specifically, stations are categorized into three performance groups[34] based on the NSE. One representative station is selected from each category, as summarized in Table 1.

Table 1 Representative stations across NSE performance categories

Performance Category	NSE Range	Station ID	NSE
Very Good	$\text{NSE} > 0.75$	09066300	0.9905
Good	$0.65 < \text{NSE} \leq 0.75$	01439500	0.7366
Satisfactory	$0.50 < \text{NSE} \leq 0.65$	03504000	0.6217

We first consider a representative inland station with very good performance (Station ID: 09066300, *Middle Creek near Minturn, CO.*, 39.65°N, 106.38°W, NSE = 0.9905) to illustrate the extracted core circuit (Fig. 1a). The station is located in a high-altitude inland basin in Colorado, USA (approximately 3251 m), where hydrological processes are strongly influenced by temperature-driven dynamics and catchment routing[35].

The circuit structure (Fig. 1a) shows that the model retains a clear and stable information flow pathway after strong sparsification. Among the inputs, streamflow (Q) and maximum temperature (Tmax) dominate, with substantially stronger connections than DayL. Beyond variable importance, this structure suggests a process-level pattern

in which antecedent catchment state (with Q acting as a storage proxy) and temperature jointly regulate runoff release. The recovered pathway is consistent with plausible high-altitude controls such as snowmelt timing and evapotranspiration sensitivity[36]. Within the hidden layers, information propagates through a small number of key nodes and converges to the output, forming a compact, low-redundancy computational pathway.

Training dynamics (Fig. 1d, lower right) provide complementary evidence: as the number of retained nodes (K) increases, NSE rapidly improves to near 1 and stabilizes when K remains relatively small ($K \leq 100$), while MSE decreases sharply and stays low. Notably, performance comparable to the full model is achieved with only a small subset of nodes. This indicates that the extracted sparse circuit preserves essential predictive capability, supporting its functional fidelity[37]. More broadly, the sparsity-inducing training appears to retain critical computational units while removing redundant structure[38].

Furthermore, the interaction analysis (Fig. 1e) reveals strong dependencies among input variables. In particular, T_{max} exhibits the strongest interaction with Q , followed by the interaction between Q and $DayL$, while $DayL$ shows relatively weak interactions overall. This supports a process-level interpretation: thermal forcing does not act independently, but modulates discharge response through antecedent-state conditioning, consistent with energy-state coupling in mountain basins[39].

Finally, the contribution analysis (Fig. 1f) shows a clear imbalance among input variables. Streamflow (Q) contributes the most, followed by T_{max} , while $DayL$ has the smallest contribution. Combined with the circuit and interaction evidence, this indicates that the dominant hydrological pattern at this site is a state-dependent runoff pathway with secondary temperature regulation, rather than a uniformly distributed multi-driver response.

Overall, for this very good inland station, the extracted core circuit is highly sparse while maintaining strong predictive capability[40]. More importantly, STEH recovers a coherent hydrological mode in which antecedent state and thermal forcing jointly govern streamflow response, suggesting that the learned sparse circuit reflects hydrologically meaningful structure rather than only statistical attribution.

To test whether this pattern persists beyond the best case, we next examine a representative station with good performance (Station ID: 01439500, *Bush Kill at Shoemakers, PA*, 41.09°N, 75.04°W, NSE = 0.7366). The extracted circuit exhibits a moderately sparse structure (Fig. 1b). Compared with the very good inland station, a broader set of inputs contribute to prediction, including precipitation (P), streamflow (Q), and radiation-related variables. The resulting information flow is therefore more distributed[38], consistent with multiple sub-processes acting simultaneously. Interaction analysis still shows a dominant role for Q , but with stronger coupling to precipitation, which aligns with an event-response pathway where rainfall signals are filtered by antecedent wetness before being transmitted to discharge[41].




Fig. 1 Sparse Circuit Extraction and Interpretability Analysis of Hydrological Forecasting Models(a), (b), (c) Core Circuit Structures: Deep learning model circuits for high-performance (NSE = 0.9905), medium-performance (NSE = 0.7366), and low-performance (NSE = 0.6217) sites, showing sparse connectivity from input to output layers.(d) Input Variable Interaction Heatmap: Pairwise dependencies between input variables, with darker colors indicating stronger interactions.(e) Node Pruning Training Curves: Performance and error evolution as node count changes, revealing model redundancy and extraction effectiveness.(f) Input Variable Contribution Diagram: Visualization of input variable contributions to flow prediction, with chord width indicating influence magnitude.*In panel (d), performance decreases from left to right; in panels (e) and (f), top to bottom represents performance from optimal to worst.*

When performance decreases further, the learned structure becomes visibly less coherent. For a representative station with satisfactory performance (Station ID: 03504000, *Nantahala River near Rainbow Springs, NC*, 35.13°N, 83.62°W, NSE = 0.6217), the extracted circuit is less structured and more diffuse (Fig. 1c). Although streamflow (Q) still dominates prediction, other variables such as radiation (SRAD) and DayL show weaker and less consistent contributions. Interaction patterns are also less pronounced, suggesting that no single robust hydrological pathway is cleanly separated from competing signals. The circuit retains more redundant connections, and sparsity-inducing training appears less effective in isolating a compact predictive core[36].

In our comparative analysis, we observed an interesting phenomenon: as prediction performance decreases, the sparsity of the extracted circuits also decreases, and the structure tends to become more dispersed. Specifically, high-performance sites exhibit compact and orderly pathways, which correspond to clear hydrological patterns, while low-performance sites show increased redundancy, weaker interaction consistency, and greater instability in the interpretability of the knowledge layer[42]. This trend has attracted our attention. As a result, we conducted a more systematic analysis to verify *whether there is a significant and stable relationship between the degree of model structure simplification and prediction performance*.

2.2 Parsimony and Physical Interpretability Are Prevalent

Interestingly, our findings reveal a trend that contrasts with observations reported in some prior studies[43, 44], where increased model complexity has been associated with improved predictive performance. As shown in Table 2, models with higher predictive skill ($\text{NSE} \geq 0.5$) exhibit significantly lower structural complexity compared to lower-performing counterparts, as indicated by a reduced number of nodes (45.77 vs. 56.35, $p = 5.00 \times 10^{-5}$). This suggests that increased model complexity does not necessarily translate into improved performance in this context; instead, more parsimonious architectures may facilitate better generalization[42]. One possible explanation is that overly complex models introduce redundant representations or noise, which can degrade predictive robustness in hydrological applications[45].

Table 2 Structural complexity across NSE-based performance groups

Group	Stations	Mean Nodes	Median Nodes
$\text{NSE} < 0.5$	194	56.35	56.50
$\text{NSE} \geq 0.5$	327	45.77	42.00
All	521	49.71	49.00

Note: A statistically significant difference is observed between groups ($p = 5.00 \times 10^{-5}$).

From the perspective of physical interpretability, the probe analysis (Table 3) indicates that key hydrological variables—including CUM, API, and CHG—are consistently encoded in the learned representations across models[46]. Specifically, no statistically significant differences are observed in probe performance between NSE-based groups

(all $p > 0.05$). This implies that the ability to capture physically meaningful signals is broadly preserved across model groups[22], rather than being exclusive to high-performing models.

Table 3 Probe performance under different NSE conditions

Probe	Group	Mean	Median	p-value
R_{cum}^2	NSE < 0.5	0.8746	0.9639	0.4341
	NSE $\geq$ 0.5	0.8915	0.9617	
	All	0.8851	0.9621	
R_{api}^2	NSE < 0.5	0.8795	0.9682	0.5676
	NSE $\geq$ 0.5	0.8917	0.9651	
	All	0.8871	0.9663	
R_{chg}^2	NSE < 0.5	0.3791	0.4267	0.1563
	NSE $\geq$ 0.5	0.4054	0.4406	
	All	0.3955	0.4342	

Outliers were excluded using the 1.5 IQR criterion within each group. Results remain consistent without filtering.

Taken together, these results suggest that parsimony and physical interpretability are not competing objectives. Models with simpler structures can still achieve strong predictive performance while encoding essential hydrological knowledge related to streamflow generation. This naturally motivates the next question: *if such knowledge is broadly present across models, what is its concrete content and spatial organization across basins?*

2.3 Core Circuits Reflect Hydrological Knowledge Across Basins

To answer that question, we examine the spatial organization of the learned interaction structures within core circuits. The results indicate that the model captures structured and regime-dependent hydrological knowledge across diverse basins[47]. As shown in Fig. 2(a), the dominant interaction in most basins is between precipitation and discharge (P–Q). This pattern reflects the fundamental role of precipitation as the primary driver of streamflow generation[48]. The detailed procedure for how we performed the classification can be found in the Supplementary Information 3.1.

In contrast, a subset of basins is characterized by strong interactions between lagged discharge and target discharge (Q–Q_{target}), suggesting pronounced memory effects[4]. In these basins, streamflow dynamics are more strongly influenced by antecedent conditions than by immediate precipitation inputs, which is typically associated with limited precipitation or strong storage capacity. This pattern can be interpreted as a storage-controlled response mode, where antecedent wetness and catchment retention regulate subsequent flow release.

Temperature-related interactions (T_{max}–Q and T_{min}–Q) are primarily observed in inland and high-elevation regions. These patterns indicate that temperature plays a




Fig. 2 Spatial patterns of hydrological interaction regimes and their performance (a,c,e) Spatial distribution of dominant interaction modes (e.g., P–Q, Q–Q_{target}, and energy-related). (b,d,f) NSE, interaction strength, and physical probe states for the corresponding interaction types.




Fig. 3 Distribution of dominant interaction groups across different hydro-climatic regimes, including aridity levels, snow influence, and geo-climatic types.

critical role in regulating streamflow generation in such environments, not only through snowmelt processes but also via evaporation and other energy-limited controls[3]. In addition, interactions involving shortwave radiation (SRD) emerge in some basins, further highlighting the importance of energy inputs[2] in modulating hydrological responses. Together, these signals correspond to an energy-constrained runoff mode, distinct from precipitation-dominated systems.

These patterns remain consistent across different hydro-climatic conditions, as illustrated in Fig. 2(b–e). Precipitation-driven interactions dominate across most basins, whereas temperature-related and lagged discharge interactions become more prominent under specific environmental settings. In arid and semi-humid regions, temperature and memory-related interactions are more frequently observed[3], indicating stronger energy constraints and persistence effects. In contrast, humid basins are largely governed by precipitation-driven dynamics[48].

Similarly, temperature-related interactions[2] are concentrated in snow-influenced basins, consistent with snowmelt-controlled streamflow processes, while precipitation-driven interactions dominate in rain-fed systems[49]. Across different geo-climatic types, distinct interaction structures exhibit clear regional preferences, suggesting that the learned circuit structures adapt to local environmental conditions.

Overall, these findings indicate that the learned core circuits are not arbitrary but systematically reflect regime-dependent hydrological patterns[50]. Thus, STEH moves the interpretation target from “which variable is important” to “which process pathway is operating,” providing structural evidence that the model internalizes physically meaningful knowledge beyond simple variable associations. However, structural coherence alone is insufficient: the extracted knowledge should also produce physically consistent behavior under controlled perturbations.

2.4 Accurate Volume Response but Limited Sensitivity to Temporal and State Perturbations

Following the structural analysis above, we next test whether the extracted hydrological knowledge corresponds to realistic model behavior. As shown in Fig. 4, the counterfactual analysis indicates that the model exhibits physically meaningful responses under different types of perturbations[31]. For magnitude-based interventions, particularly rainfall scaling (Fig. 4b), the model produces directionally consistent and approximately linear changes in streamflow volume across stations. As precipitation increases (or decreases), the predicted streamflow volume correspondingly increases (or decreases), while prediction error remains largely unchanged. This indicates that the model has learned a stable and physically consistent mapping between precipitation and streamflow generation, and can respond to changes in input water magnitude without substantial loss of predictive accuracy[28].



Fig. 4 Counterfactual sensitivity analysis across intervention types. (a) Station-level expected behavior statistics. (b) Rainfall scaling. (c) Rainfall pattern redistribution. (d) streamflow-state scaling. (e) Temperature perturbation. Blue lines denote relative changes in predicted streamflow volume, and orange lines denote relative changes in prediction error. Detailed experimental design and evaluation metrics are provided in Supplementary Information 3.2.

Under rainfall pattern redistribution (Fig. 4c), the total precipitation is conserved, and the predicted streamflow volume remains broadly stable, consistent with mass conservation expectations[51]. Structural differences nevertheless emerge across redistribution modes. In particular, the back-loaded configuration shows a clear positive deviation in predicted streamflow volume relative to other scenarios. This indicates that the model is sensitive to the temporal concentration of rainfall and produces stronger streamflow responses[47] when precipitation is concentrated toward later periods, reflecting an initial capability to capture rainfall-streamflow lag relationships. Meanwhile, the variation in prediction error across scenarios indicates that the model’s utilization of temporal structure remains somewhat scenario-dependent.

For streamflow-state perturbations (Fig. 4d), moderate scaling of recent streamflow history leads to corresponding adjustments in predicted volume, while prediction error remains stable. This indicates that the model can utilize state information while maintaining consistent predictive performance. As the scaling magnitude increases, however, the model response gradually saturates, with volume changes no longer increasing proportionally to the input. This attenuation pattern suggests that the model does not simply amplify streamflow magnitude linearly, but may impose constraints on non-structured or inconsistent input variations, thereby avoiding unrealistic over-amplification. Such behavior is consistent with a model that captures streamflow dynamics through interactions among multiple hydrological factors[52] rather than relying solely on absolute magnitude.

For temperature perturbations (Fig. 4e), the predicted streamflow volume shows a slight decreasing trend, with consistent direction across snow-dominated and rainfall-dominated catchments and no significant difference between groups. This aligns with hydrological expectations that warming generally reduces streamflow through enhanced evapotranspiration or phase-related processes. In contrast, prediction error exhibits different sensitivities between the two regimes: snow-dominated catchments show substantially stronger responses to temperature changes than rainfall-dominated ones[3], and this difference is consistent across warming magnitudes ($p < 0.001$). This indicates that the model not only captures the overall volumetric effects of temperature, but also differentiates between distinct hydrological knowledge patterns, reflecting sensitivity to process heterogeneity.

Overall, the model demonstrates robust responses to precipitation-driven volume changes and selective, generally limited sensitivity to temporal structure and catchment state variations[53], indicating effective learning of key hydrological knowledge with remaining gaps in temporal-state representation. At the same time, the differing response patterns across perturbation types suggest that the representation of more complex processes, such as temporal organization and state evolution, remains an area for further improvement. This behavioral consistency supports the validity of the extracted knowledge, and naturally leads to the final question: *can this knowledge be transferred to improve other model classes?*

2.5 Cross-Model Transfer Validates STEH-Derived Knowledge

To answer this final question, we evaluate transferability using an external model family. Specifically, we designed a comparative experiment in which STEH-derived structural

information is incorporated into a conventional Random Forest model[54] and compared against a baseline model without such information. The key purpose is not only to improve another model, but to test whether the extracted structures reflect model-independent hydrological knowledge that remains effective outside the Transformer architecture[55].

Furthermore, station-wise paired tests (paired t-test and Wilcoxon signed-rank test) reported in Supplementary Information 3.3 indicate that the modified model outperforms the baseline with statistical significance, supporting the robustness of the observed improvements.

From an overall perspective, the modified model exhibits consistent performance improvements across stations[56]. Specifically, NSE increases from 0.760 to 0.767, accompanied by reductions in both MSE and MAE, indicating improved predictive accuracy and error control[57]. The station-wise comparison (Fig. 5(d)) shows that most points lie above the diagonal line, and the gains persist across a wide range of baseline performance levels (Fig. 5(e)). These observations indicate that the transferred structures retain predictive value in an independent model class, consistent with the claim that STEH extracts hydrological knowledge beyond architecture-specific heuristics[58], while the magnitude of improvement remains modest at the aggregate level.

Table 4 Performance comparison between baseline and STEH-informed modified Random Forest models across hydro-climatic groups

Model	Group	MSE	MAE	NSE
Baseline	Arid	0.140054	0.260355	0.841540
	Humid	0.324556	0.435797	0.700072
	Sub-humid	0.270571	0.385097	0.724154
	All Stations	0.237974	0.353715	0.759683
Modified	Arid	0.133525	0.248404	0.846403
	Humid	0.313798	0.425727	0.709464
	Sub-humid	0.259417	0.372654	0.733143
	All Stations	0.228561	0.342049	0.767310

All metrics are averaged across stations within each hydro-climatic group. Only stations with STEH model performance (NSE > 0.5) are included, as the extracted hydrological structures are derived from the STEH model and may become unreliable when its predictive performance is low.

Across different hydro-climatic regimes, the most pronounced improvements are observed in arid regions. As shown in Fig. 5(a) and Table 4, inland arid and semi-arid areas exhibit predominantly positive performance gains. In the Arid group, the modified model consistently outperforms the baseline in terms of MSE, MAE, and NSE, suggesting that incorporating structural information can enhance predictive capability under such conditions. This is consistent with the intuition that in environments where precipitation signals are weak and hydrological responses are governed by multiple interacting factors[59], additional structural features can help the model better capture streamflow dynamics.

In humid regions, the modified model also achieves consistent improvements (NSE:




Fig. 5 Performance comparison between baseline and STEH-informed Random Forest models. (a) Spatial distribution of NSE differences (red: improved; blue: degraded). (b,c) Performance distributions of the baseline and modified models. (d) Station-wise NSE comparison (points above the diagonal indicate improvement). (e) NSE improvement relative to baseline performance. All results are based on stations with STEH model performance ($NSE > 0.5$).

0.700 to 0.709), although the magnitude of improvement is relatively smaller. This can be attributed to the comparatively simpler hydrological processes in such regions[60], where streamflow generation is primarily controlled by a dominant driver. Under these conditions, the baseline model is already capable of capturing the main dynamics, leaving limited room for further improvement. Nevertheless, the continued reduction in error metrics indicates that structural information still contributes to improved fine-scale representation[61].

Overall, these findings demonstrate that the structural information extracted by STEH is not only interpretable but also scientifically verifiable through cross-model transfer. In particular, in regions characterized by complex or multi-factor-controlled processes, incorporating such structural knowledge can substantially improve model performance. Successful transfer to Random Forest provides an external validation step: the extracted knowledge is not tied to a single network design, but reflects portable hydrological regularities. In this section we successfully illustrate several major questions: hydrological knowledge is broadly encoded across models, can be structurally identified, can be behaviorally validated for physical consistency, and can be externally verified through cross-model transfer.

2.6 Limitations

Nevertheless, despite the progress made in uncovering knowledge representations in deep learning-based hydrological models, several limitations should be acknowledged.

First, although STEH improves structural interpretability, its results still cannot be interpreted as strict causal relationships [62]. This limitation arises from the intrinsic nature of deep learning models as statistical learning systems [63], where extracted patterns primarily reflect statistical associations rather than true physical causality. Therefore, the findings should be interpreted with caution when relating them to real-world physical processes.

Second, although this study reveals interaction structures among variables, their theoretical linkage to explicit physical process equations remains to be further explored. For example, the specific coupling relationships between variables, the pathways through which interactions operate, and the development of more controllable and interpretable circuit structures are not yet fully characterized [64]. Moreover, although the current framework enforces sparsity to simplify internal structures, the analysis is still largely limited to the relationship between input variables and the final hidden representations, leaving the information evolution within intermediate high-dimensional feature spaces insufficiently explored.

In addition, the validation of the extracted “physical knowledge” in this study is still based on machine learning models rather than explicit process-based hydrological models [65]. This may limit the physical rigor of the conclusions, and future work should incorporate process-driven models to further verify the findings.

3 Summary

Deep learning has advanced hydrological prediction, but model interpretability remains challenged by low verifiability, neglected variable dependencies[22], and unclear internal structures[66], leaving uncertain whether high-performance models learn physically meaningful hydrological processes[67]. This study proposes the Sparse Transformer-based Explainable Hydrology (STEH) framework. It establishes an end-to-end interpretability pipeline using structured sparsification, circuit-level ablation, variable interaction analysis[38], and representation probing to systematically reveal the internal decision-making mechanisms of deep learning models in hydrology.

The main findings are as follows. First, STEH provides a circuit-level view to explicitly identify decision pathways and information flow, offering a structurally grounded and verifiable interpretability paradigm. Second, simpler models form more compact internal circuits, while over-complex models introduce redundant structures; model parsimony improves clarity while preserving physical meaning. Third, physically meaningful representations are learned across models regardless of performance, indicating physical knowledge is inherent in data-driven models. Fourth, STEH reliably characterizes stable and hydrologically consistent variable interactions across diverse hydroclimatic regimes. Fifth, sensitivity analysis shows the model effectively captures dominant streamflow-generation knowledge but incompletely represents temporal organization and state evolution. Sixth, STEH-derived knowledge is transferable and practically useful, as validated by performance improvements in a Random Forest model. Seventh, this study provides a replicable pathway for interpretability and physical validation of deep learning models in hydrological research.

In summary, this study demonstrates that hydrological deep learning models encode physically meaningful information that can be systematically uncovered through structural analysis. The STEH framework bridges black-box prediction and hydrological interpretation, supports knowledge discovery, and offers a generalizable route for integrating data-driven and physics-based hydrological modeling.

4 Methods and data

4.1 Framework Overview

The overall workflow of the proposed Sparse Transformer-based Explainable Hydrology (STEH) framework is illustrated in Fig. 6. The framework provides an end-to-end interpretability pipeline for deep learning-based streamflow prediction, enabling systematic analysis from hydrological input variables, through internal circuit structures, to prediction outputs.

The STEH framework consists of three main components: (1) sparse circuit discovery, (2) circuit verification and refinement, and (3) physical interpretation. Together, these components aim to reveal how deep learning models process hydrological information and whether the learned representations correspond to meaningful hydrological processes. Terminology follows the Appendix 1 terminology table (Table 5).

First, sparse circuit discovery is performed by introducing learnable gating mecha-

nisms into the Transformer architecture. Gates are applied to input variables, attention heads, and feed-forward neurons, allowing the model to automatically identify a compact set of critical computational units during training. Through sparsity-inducing optimization, redundant structures are gradually suppressed, resulting in a sparse computational circuit responsible for streamflow prediction.

Second, circuit verification and refinement are conducted through systematic ablation experiments. Individual computational units—including input nodes, attention heads, and neurons—are selectively removed to measure their structural contribution to prediction performance. Units with negligible influence are discarded, yielding a compact core predictive circuit that preserves the essential functional structure of the model.

Third, the extracted core circuit is further analyzed to understand internal information flow and structural dependencies. Interaction analysis is used to identify cooperative or redundant relationships between model components. In addition, linear probing techniques are applied to hidden representations to evaluate whether internal states encode physically meaningful hydrological variables.

Finally, the interpretability results are validated from a hydrological perspective. Based on the structural knowledge extracted from the neural network, modifications are introduced to a Random Forest benchmark model to examine whether the discovered hydrological knowledge can improve data-driven hydrological simulations.

In what follows, we first introduce the sparse Transformer architecture and the three-stage sparsification strategy, and then describe the interpretability analyses and evaluation metrics.




Fig. 6 Overall architecture of the proposed STEH framework

4.2 Sparse Transformer

The Transformer architecture has recently been introduced into hydrological time-series modeling because its self-attention mechanism can effectively capture long-range temporal dependencies in hydro-meteorological sequences[68]. In streamflow prediction, antecedent rainfall, temperature, and soil moisture conditions jointly influence discharge responses across multiple time lags, making long-term dependency modeling particularly important[56]. The Transformer assigns different weights to each input feature through the self-attention mechanism, automatically learning the relationships between different time steps[69].




Fig. 7 Three-stage sparsification framework for predictive circuit discovery in Transformer. The model is trained through three stages to progressively transform a dense Transformer into a sparse predictive circuit. (a–b) Stage I: Origin model training, where the Transformer is trained without sparsity constraints and all input features contribute to the prediction. (c–d) Stage II: Weight sparsification progressively prunes redundant connections according to the sparsity schedule. (e–f) Stage III: Mask learning optimizes learnable gates so that only a subset of input features contributes to the final prediction. For clarity, only a subset of connections between layers is illustrated in the figure, while the actual model remains fully connected, and decoder attention layers are omitted for simplicity. Note that the proposed method does not explicitly modify the network architecture; sparsity is achieved by setting weights or gates to zero, while the visual removal of units in the figure is only for illustrative purposes. The number of input variables shown in the figure is fewer than the actual input feature set; this simplification is adopted solely for visual clarity.

To enhance the interpretability of the model, we introduce a structured gating mechanism to filter input variables, attention heads, and feed-forward neurons[70]. Through

these gating mechanisms, we can selectively retain input features or parts of the network structure, thus enabling more efficient computation (Fig. 7(e–f)). By controlling the activation of gates, we are able to gradually filter out the features and pathways that have the most significant impact on the prediction. This structured sparsification approach helps the model focus on the most important inputs and computations, improving computational efficiency and enhancing model interpretability[71].

To avoid prematurely removing potentially important computational pathways, our training process is divided into three stages[37].:

Stage I (Origin Model Training): The model is first trained without sparsity constraints to establish a stable predictive baseline. In this stage, all input features and computational pathways are retained to ensure the model has sufficient expressive capacity (Fig. 7(a–b)).

Stage II (Weight Sparsification): In this stage, weight sparsification is gradually imposed, reducing redundant connections in the network. We use an annealed sparsity schedule to progressively increase the sparsity ratio, controlling which connections are pruned[72]. The process sets a target sparsity ratio (from the initial $\rho_{\max}$ to the final $\rho_{\min}$) to progressively reduce the complexity of the model (S2(c–d)).

Stage III (Mask Learning): Finally, learnable gates are optimized to further refine the network (Fig. S2(e–f)). We regularize the activation probabilities of the gates, guiding the network to select the most important input features and computational pathways[73]. The goal is to make the computational pathways of the network more sparse by adjusting the gate activation probabilities, improving the model’s inference efficiency.

These stages combined result in a compact predictive circuit that provides efficient and accurate predictions with fewer computational resources[36]. Subsequently, we conduct an interpretability analysis of the model across multiple dimensions and evaluation metrics. If you are interested in the specific methods, please refer to the Supplementary Information 2 for further details (e.g., formal model definitions, interpretability techniques, etc.).

4.3 Data

This study uses the CAMELS-US dataset ([https://doi.org/10.1016/0022-1694\(89\)90184-4](https://doi.org/10.1016/0022-1694(89)90184-4)), which provides daily hydro-meteorological time series for minimally disturbed catchments across the contiguous United States[74]. We select this dataset for three reasons: (1) catchments are weakly affected by direct human regulation, which helps identify more robust hydrological patterns; (2) long and standardized daily forcing–streamflow records support deep learning time-series training and cross-basin comparison[47]; and (3) the data organization aligns with large-sample hydrology workflows, enabling reproducible evaluation[28]. To ensure temporal consistency across stations, we use a common analysis period from 1980-01-01 to 2014-12-31. Based on this period, a total of 521 stations are retained after data quality control and screening. The model input variables consist of meteorological forcing and antecedent hydrological information. Specifically, the meteorological forcing includes multiple daily atmospheric variables, while the antecedent hydrological information is represented by historical streamflow sequences to capture catchment state memory. Additional details can be found in the Supplementary Information 2.4.

The prediction task is formulated as a sequence-to-one regression problem, where a sliding window of length $L = 15$ days is used to predict next-day streamflow. All features are standardized using training-set statistics, and the data are split chronologically into training (70%), validation (15%), and test (15%) subsets.

The model is implemented in PyTorch (Python 3.10) and trained on a CUDA-enabled Linux server. A three-stage pipeline (dense pre-training, weight sparsification, and mask learning) is optimized using AdamW with weight decay. Hyperparameters are tuned via Bayesian optimization[75] with validation NSE as the objective, and the selected configuration is consistently used across all stages and subsequent explainability analyses. Further details are provided in Supplementary Information 4.1.

Data availability

The data that support the findings of this study are publicly available. The streamflow and meteorological data are derived from the CAMELS-US dataset[74]. The dataset is accessible via: [https://doi.org/10.1016/0022-1694\(89\)90184-4](https://doi.org/10.1016/0022-1694(89)90184-4).

Code availability

The source code and processed datasets are available in the GitHub repository: <https://github.com/Polyethylenekjc/Sparse-Transformer-In-Streamflow-Prediction>.

CRedit authorship contribution statement

Jiachen Kong: Conceptualization, Methodology, Software, Writing – original draft.
Zhaocai Wang: Methodology, Software, Validation, Supervision, Writing – review & editing. **Zhaoyang Zhu:** Data curation, Writing – original draft.

Declaration of Competing Interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Table 5 Appendix 1: Main Terminology Overview

Full Name	Abbreviation	Description
Physics-Informed Neural Networks	PINNs	Optimization constraint, emphasizes that predictions must obey physical laws, e.g., streamflow predictions satisfy mass conservation.
Sparse Transformer for Explainable Hydrology	STEH	Architecture design, emphasizes sparse and analyzable internal paths, e.g., identifying paths responsible for evapotranspiration.
Circuit	Circuit	Functional subset of computational units selected by sparsification (e.g., specific heads/neurons) that jointly supports prediction.
Structural	Structural	Organization pattern of components and interactions within a circuit (e.g., layer-wise arrangement and dependency layout).
Information Flow Pathway	Pathway	Directed route of information propagation within a circuit (e.g., precipitation-to-runoff transfer through key nodes).
Explainable AI	XAI	Overall goal, broad transparency and interpretability, e.g., understanding why the model predicts more accurately during drought periods.

Appendix 2: Confusable Concepts and Verification Methods

Full Name	Verification / Implementation Method
Explainability	Post-hoc analysis (SHAP / LIME) or intrinsic structural analysis to understand model prediction behavior.
Physical Structure Explainability	Structural sparsification + linear probing to verify whether internal computational paths correspond to physical quantities (e.g., soil moisture).
Causal Contribution	Rigorous causal verification, including causal graphs / SEMs, intervention experiments (Do-Calculus), and counterfactual / A/B comparisons to ensure true causal effect from input to output.
Structural Causal Contribution	Structured intervention + causal pathway verification + counterfactual comparison, to verify that internal components (attention heads, paths) are necessary and irreplaceable in the input–output causal chain.

References

- [1] Yingying Gao, Zengliang Luo, Huan Liu, Lunche Wang, Xi Chen, and Huan Li. Reconstruction of global long-term daily streamflow dataset using machine learning models for revealing streamflow changes. *Journal of Hydrology: Regional Studies*, 64:103148, 2026. doi: <https://doi.org/10.1016/j.ejrh.2026.103148>. URL <https://www.sciencedirect.com/science/article/pii/S2214581826000467>.
- [2] Walter W Immerzeel, Ludovicus PH Van Beek, and Marc FP Bierkens. Climate change will affect the asian water towers. *science*, 328(5984):1382–1385, 2010. doi: <https://doi.org/10.1126/science.1183188>. URL <https://doi.org/10.1126/science.1183188>.
- [3] Paul CD Milly and Krista A Dunne. Colorado river flow dwindles as warming-driven loss of reflective snow energizes evaporation. *Science*, 367(6483):1252–1255, 2020. doi: <https://doi.org/10.1126/science.aay9187>. URL <https://doi.org/10.1126/science.aay9187>.
- [4] Sina Apak, Huseyin Cagan Kilinc, Adem Yurtsever, Hilal Haznedar, and Furkan Ozkan. Multi-resolution adaptive channel fusion transformer encoder lstm for accurate streamflow prediction. *Scientific Reports*, 2026. doi: <https://doi.org/10.1038/s41467-026-70523-y>. URL <https://doi.org/10.1038/s41467-026-70523-y>.
- [5] Sam Zipper, Ian Gambill, Monty Schmitt, Claire Kouba, Leland Scantlebury, Thomas Harter, and Nicholas P. Murphy. Lagged streamflow depletion due to pumping-induced stream drying: Incorporation into analytical streamflow depletion estimation methods. *Journal of Hydrology*, 667:134909, 2026. doi: <https://doi.org/10.1016/j.jhydrol.2026.134909>. URL <https://www.sciencedirect.com/science/article/pii/S0022169426000065>.
- [6] Jorge Núñez, Catalina B. Cortés, and Marjorie A. Yáñez. Explainable artificial intelligence in hydrology: Interpreting black-box snowmelt-driven streamflow predictions in an arid andean basin of north-central chile. *Water*, 15(19), 2023. doi: <https://doi.org/10.3390/w15193369>. URL <https://www.mdpi.com/2073-4441/15/19/3369>.
- [7] Shuo Wang and Hui Peng. Multiple spatio-temporal scale runoff forecasting and driving mechanism exploration by k-means optimized xgboost and shap. *Journal of Hydrology*, 630:130650, 2024. doi: <https://doi.org/10.1016/j.jhydrol.2024.130650>. URL <https://www.sciencedirect.com/science/article/pii/S0022169424000441>.
- [8] Jacob Mardian, Catherine Champagne, Barrie Bonsal, and Aaron Berg. A machine learning framework for predicting and understanding the canadian drought monitor. *Water Resources Research*, 59(8):e2022WR033847, 2023. doi: <https://doi.org/10.1029/2022WR033847>. URL <https://agupubs.onlinelibrary.wiley.com/doi/abs/10.1029/2022WR033847>.

- [9] Felix Schmid, Leonie Müller, and Jorge Leandro. A physically informed domain-independent data-driven inundation forecast model. *Water Research*, 289:124819, 2026. doi: <https://doi.org/10.1016/j.watres.2025.124819>. URL <https://www.sciencedirect.com/science/article/pii/S0043135425017221>.
- [10] Zheng Fang, Simin Qu, Ziheng Li, Qiongfang Li, Peng Shi, Yiqun Sun, Xiaoliang Yang, Han Tang, Junfang Zhang, Zhengxian Zhu, Hongshi Wu, and Weimin Bao. Soil water accounting network (swan): a novel neural network for modeling conceptual hydrological processes. *Journal of Hydrology*, 661:133562, 2025. doi: <https://doi.org/10.1016/j.jhydrol.2025.133562>. URL <https://www.sciencedirect.com/science/article/pii/S002216942500900X>.
- [11] Mingrui Shi, Hongyuan Fang, Yangyang Xie, Huihua Du, Saiyan Liu, Jean Marie Ndayiragije, and Nannan Liu. A flood forecasting method coupling the ce-qual-w2 and pinn models. *Journal of Hydrology*, 661:133699, 2025. doi: <https://doi.org/10.1016/j.jhydrol.2025.133699>. URL <https://www.sciencedirect.com/science/article/pii/S0022169425010376>.
- [12] Wenli Liu, Qian Wu, Zihan Liu, Xu Zheng, Tianxiang Liu, and Han Gao. Pinn framework for urban flood depth prediction: integrating data-driven insights with physical constraints. *Journal of Hydrology*, 665:134693, 2026. doi: <https://doi.org/10.1016/j.jhydrol.2025.134693>. URL <https://www.sciencedirect.com/science/article/pii/S0022169425020335>.
- [13] Ribana Roscher, Bastian Bohn, Marco F. Duarte, and Jochen Garcke. Explainable machine learning for scientific insights and discoveries. *IEEE Access*, 8:42200–42216, 2020. doi: <https://doi.org/10.1109/ACCESS.2020.2976199>. URL <https://doi.org/10.1109/ACCESS.2020.2976199>.
- [14] M.Z. Naser. Fundamental flaws of physics-informed neural networks and explainability methods in engineering systems. *Computers & Industrial Engineering*, 212:111704, 2026. doi: <https://doi.org/10.1016/j.cie.2025.111704>. URL <https://www.sciencedirect.com/science/article/pii/S0360835225008502>.
- [15] Aditi Krishnapriyan, Amir Gholami, Shandian Zhe, Robert Kirby, and Michael W Mahoney. Characterizing possible failure modes in physics-informed neural networks. In *Advances in Neural Information Processing Systems*, volume 34, pages 26548–26560, 2021. doi: <https://dl.acm.org/doi/10.5555/3540261.3542294>. URL https://proceedings.neurips.cc/paper_files/paper/2021/file/df438e5206f31600e6ae4af72f2725f1-Paper.pdf.
- [16] Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. ”why should i trust you?”: Explaining the predictions of any classifier. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, page 1135–1144, 2016. doi: <https://doi.org/10.1145/2939672.2939778>. URL <https://doi.org/10.1145/2939672.2939778>.

- [17] Scott Lundberg and Su-In Lee. A unified approach to interpreting model predictions, 2017. URL <https://arxiv.org/abs/1705.07874>.
- [18] Stephanie R. Clark, Guobin Fu, and Sreekanth Janardhanan. Explainable ai for interpreting spatiotemporal groundwater predictions. *Water Resources Research*, 61(10):e2025WR041303, 2025. doi: <https://doi.org/10.1029/2025WR041303>. URL <https://agupubs.onlinelibrary.wiley.com/doi/abs/10.1029/2025WR041303>.
- [19] Shah Mehmood Wagan and Sidra Sidra. Interpreting irrigation decisions: Explainable ai using shap and lime in agricultural water management models. *Water Resources Management*, 40(52), 2026. doi: <https://doi.org/10.1007/s11269-025-04468-0>. URL <https://doi.org/10.1007/s11269-025-04468-0>.
- [20] Indra Kumar, Carlos Scheidegger, Suresh Venkatasubramanian, and Sorelle Friedler. Shapley residuals: Quantifying the limits of the shapley value for explanations. In *Advances in Neural Information Processing Systems*, volume 34, pages 26598–26608, 2021. doi: <https://dl.acm.org/doi/10.5555/3540261.3542298>. URL https://proceedings.neurips.cc/paper_files/paper/2021/file/dfc6aa246e88ab3e32caaaaecf433550-Paper.pdf.
- [21] Yoshiyasu Takefuji. Limitations of shap-based interpretations in environmental and membrane filtration applications. *Water Research*, 288:124766, 2026. doi: <https://doi.org/10.1016/j.watres.2025.124766>. URL <https://www.sciencedirect.com/science/article/pii/S0043135425016690>.
- [22] Afek Ilay Adler and Amichai Painsky. Feature importance in gradient boosting trees with cross-validation feature selection. *Entropy*, 24(5), 2022. doi: <https://doi.org/10.3390/e24050687>. URL <https://www.mdpi.com/1099-4300/24/5/687>.
- [23] Andrew J. Newman, Naoki Mizukami, Martyn P. Clark, Andrew W. Wood, Bart Nijssen, and Grey Nearing. Benchmarking of a physically based hydrologic model. *Journal of Hydrometeorology*, 18(8):2215 – 2225, 2017. doi: <https://doi.org/10.1175/JHM-D-16-0284.1>. URL https://journals.ametsoc.org/view/journals/hydr/18/8/jhm-d-16-0284_1.xml.
- [24] Wanjie Zhou, Shanhu Jiang, Miao He, Ruoqi Li, Mingyan Wu, and Liliang Ren. A hybrid swat-lstm model for streamflow simulation with shap-based interpretability: Application in the wei river basin, china. *Journal of Hydrology: Regional Studies*, 64:103254, 2026. ISSN 2214-5818. doi: <https://doi.org/10.1016/j.ejrh.2026.103254>. URL <https://www.sciencedirect.com/science/article/pii/S2214581826001527>.
- [25] Jinxi Xia, Ting Liu, Hao Hou, Haolong Zhang, and Tangao Hu. Urban pluvial flood susceptibility assessment for the central lin’an district in hangzhou, china,

- using the gnnwlr model with shap interpretability. *Journal of Hydrology: Regional Studies*, 64:103232, 2026. ISSN 2214-5818. doi: <https://doi.org/10.1016/j.ejrh.2026.103232>. URL <https://www.sciencedirect.com/science/article/pii/S2214581826001308>.
- [26] Manjula Devak and C. T. Dhanya. Sensitivity analysis of hydrological models: review and way forward. *Journal of Water and Climate Change*, 8(4):557–575, 06 2017. ISSN 2040-2244. doi: 10.2166/wcc.2017.149. URL <https://doi.org/10.2166/wcc.2017.149>.
- [27] M. Jeung, Y. Her, S.-S. Baek, and K. Yoon. Sensitivity of hydrological machine learning prediction accuracy to information quantity and quality. *Hydrology and Earth System Sciences*, 30(4):1077–1095, 2026. doi: 10.5194/hess-30-1077-2026. URL <https://hess.copernicus.org/articles/30/1077/2026/>.
- [28] Grey S Nearing, Frederik Kratzert, Alden Keefe Sampson, Craig S Pelissier, Daniel Klotz, Jonathan M Frame, Cristina Prieto, and Hoshin V Gupta. What role does hydrological science play in the age of machine learning? *Water Resources Research*, 57(3):e2020WR028091, 2021. doi: <https://doi.org/10.1029/2020WR028091>. URL <https://doi.org/10.1029/2020WR028091>.
- [29] Filip Simic et al. A comprehensive analysis of perturbation methods in explainable ai. *Scientific Reports*, 15(1):9538, 2025. doi: 10.1038/s41598-025-09538-2.
- [30] Yituo Zhang, Jihong Wang, Chaolin Li, Hengpan Duan, and Wenhui Wang. Attention-based deep learning models for predicting anomalous shock of wastewater treatment plants. *Water Research*, 275:123192, 2025. ISSN 0043-1354. doi: <https://doi.org/10.1016/j.watres.2025.123192>. URL <https://www.sciencedirect.com/science/article/pii/S004313542500106X>.
- [31] Jordan S Read, Xiaowei Jia, Jared Willard, Alison P Appling, Jacob A Zwart, Samantha K Oliver, Anuj Karpatne, Gretchen JA Hansen, Paul C Hanson, William Watkins, et al. Process-guided deep learning predictions of lake water temperature. *Water Resources Research*, 55(11):9173–9190, 2019. doi: <https://doi.org/10.1029/2019WR024922>. URL <https://doi.org/10.1029/2019WR024922>.
- [32] Guanchu Wang, Yu-Neng Chuang, Fan Yang, Mengnan Du, Chia-Yuan Chang, Shaochen Zhong, Zirui Liu, Zhaozhuo Xu, Kaixiong Zhou, Xuanning Cai, and Xia Hu. TVE: Learning meta-attribution for transferable vision explainer. In Ruslan Salakhutdinov, Zico Kolter, Katherine Heller, Adrian Weller, Nuria Oliver, Jonathan Scarlett, and Felix Berkenkamp, editors, *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 50248–50267. PMLR, 21–27 Jul 2024. URL <https://proceedings.mlr.press/v235/wang24j.html>.
- [33] Xinyu Li et al. Transferability and interpretability of the sepsis prediction models in the intensive care unit. *BMC Medical Informatics and Decision Making*, 22(1): 1–13, 2022. doi: 10.1186/s12911-022-02090-3.

- [34] Daniel N Moriasi, Jeffrey G Arnold, Michael W Van Liew, Ronald L Bingner, Daren Harmel, and Thomas L Veith. Model evaluation guidelines for systematic quantification of accuracy in watershed simulations. *Transactions of the ASABE*, 50(3):885–900, 2007. doi: <https://doi.org/10.13031/2013.23153>. URL <https://doi.org/10.13031/2013.23153>.
- [35] H. Biemans, C. Siderius, A. F. Lutz, S. Nepal, B. Ahmad, T. Hassan, W. von Bloh, R. R. Wijngaard, P. Wester, A. B. Shrestha, and W. W. Immerzeel. Importance of snow and glacier meltwater for agriculture on the indo-gangetic plain. *Nature Sustainability*, 2(7):594–601, 2019. doi: <https://doi.org/10.1038/s41893-019-0305-3>. URL <https://doi.org/10.1038/s41893-019-0305-3>.
- [36] Leo Gao, Achyuta Rajaram, Jacob Coxon, Soham V. Govande, Bowen Baker, and Dan Mossing. Weight-sparse transformers have interpretable circuits, 2025. URL <https://arxiv.org/abs/2511.13653>.
- [37] Song Han, Jeff Pool, John Tran, and William Dally. Learning both weights and connections for efficient neural network. In *Advances in Neural Information Processing Systems*, volume 28, 2015. doi: <https://dl.acm.org/doi/10.5555/2969239.2969366>. URL https://proceedings.neurips.cc/paper_files/paper/2015/file/ae0eb3eed39d2bcef4622b2499a05fe6-Paper.pdf.
- [38] Kevin Wang, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt. Interpretability in the wild: a circuit for indirect object identification in gpt-2 small, 2022. URL <https://arxiv.org/abs/2211.00593>.
- [39] Bingbing Ding, Xinxiao Yu, and Guodong Jia. Exploring the controlling factors of watershed streamflow variability using hydrological and machine learning models. *Water Resources Research*, 61(5):e2024WR039734, 2025. doi: <https://doi.org/10.1029/2024WR039734>. URL <https://agupubs.onlinelibrary.wiley.com/doi/abs/10.1029/2024WR039734>.
- [40] Wojciech Samek, Alexander Binder, Grégoire Montavon, Sebastian Lapuschkin, and Klaus-Robert Müller. Evaluating the visualization of what a deep neural network has learned. *IEEE Transactions on Neural Networks and Learning Systems*, 28(11):2660–2673, 2017. doi: <https://doi.org/10.1109/TNNLS.2016.2599820>. URL <https://doi.org/10.1109/TNNLS.2016.2599820>.
- [41] Donald H. Burn and Paul H. Whitfield. Climate related changes to flood regimes show an increasing rainfall influence. *Journal of Hydrology*, 617:129075, 2023. doi: <https://doi.org/10.1016/j.jhydrol.2023.129075>. URL <https://www.sciencedirect.com/science/article/pii/S0022169423000173>.
- [42] Rene Orth, Maria Staudinger, Sonia I. Seneviratne, Jan Seibert, and Massimiliano Zappa. Does model performance improve with complexity? a case study with three hydrological models. *Journal of Hydrology*, 523:147–159, 2015. doi: <https://doi.org/10.1016/j.jhydrol.2015.05.015>.

- org/10.1016/j.jhydrol.2015.01.044. URL <https://www.sciencedirect.com/science/article/pii/S002216941500061X>.
- [43] Zhiqiang Dong, Hongchang Hu, Hui Liu, Baoligao Baiyin, Xiangpeng Mu, Jie Wen, Dengfeng Liu, Lajiao Chen, Guanghui Ming, Xue Chen, and Xiaochen Li. Superior performance of hybrid model in ungauged basins for real-time hourly water level forecasting – a case study on the Lancang-Mekong mainstream. *Journal of Hydrology*, 633:130941, 2024. doi: <https://doi.org/10.1016/j.jhydrol.2024.130941>. URL <https://www.sciencedirect.com/science/article/pii/S0022169424003354>.
 - [44] Renjie Zhou, Quanrong Wang, Aohan Jin, Wenguang Shi, and Shiqi Liu. Interpretable multi-step hybrid deep learning model for karst spring discharge prediction: Integrating temporal fusion transformers with ensemble empirical mode decomposition. *Journal of Hydrology*, 645:132235, 2024. doi: <https://doi.org/10.1016/j.jhydrol.2024.132235>. URL <https://www.sciencedirect.com/science/article/pii/S0022169424016317>.
 - [45] Jordan S. Read, Xiaowei Jia, Jared Willard, Alison P. Appling, Jacob A. Zwart, Samantha K. Oliver, Anuj Karpatne, Gretchen J. A. Hansen, Paul C. Hanson, William Watkins, Michael Steinbach, and Vipin Kumar. Process-guided deep learning predictions of lake water temperature. *Water Resources Research*, 55(11):9173–9190, 2019. doi: <https://doi.org/10.1029/2019WR024922>. URL <https://agupubs.onlinelibrary.wiley.com/doi/abs/10.1029/2019WR024922>.
 - [46] Vikas Verma, Alex Lamb, Christopher Beckham, Amir Najafi, Ioannis Mitliagkas, David Lopez-Paz, and Yoshua Bengio. Manifold mixup: Better representations by interpolating hidden states. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97, pages 6438–6447, 2019. doi: <https://doi.org/10.48550/arXiv.1806.05236>. URL <https://proceedings.mlr.press/v97/verma19a.html>.
 - [47] F. Kratzert, D. Klotz, G. Shalev, G. Klambauer, S. Hochreiter, and G. Nearing. Towards learning universal, regional, and local hydrological behaviors via machine learning applied to large-sample datasets. *Hydrology and Earth System Sciences*, 23(12):5089–5110, 2019. doi: <https://doi.org/10.5194/hess-23-5089-2019>. URL <https://hess.copernicus.org/articles/23/5089/2019/>.
 - [48] W. R. Berghuijs, R. A. Woods, and M. Hrachowitz. A precipitation shift from snow towards rain leads to a decrease in streamflow. *Nature Climate Change*, 4(7):583–586, 2014. doi: <https://doi.org/10.1038/nclimate2246>. URL <https://doi.org/10.1038/nclimate2246>.
 - [49] Samuel E. Franzen, Mozghan A. Farahani, and Allison E. Goodwell. Information flows: Characterizing precipitation-streamflow dependencies in the Colorado headwaters with an information theory approach. *Water Resources Research*, 56(10):e2019WR026133, 2020. doi: <https://doi.org/10.1029/2019WR026133>.

- URL <https://agupubs.onlinelibrary.wiley.com/doi/abs/10.1029/2019WR026133>.
- [50] F. U. Jehn, K. Bestian, L. Breuer, P. Kraft, and T. Houska. Using hydrological and climatic catchment clusters to explore drivers of catchment behavior. *Hydrology and Earth System Sciences*, 24(3):1081–1100, 2020. doi: <https://doi.org/10.5194/hess-24-1081-2020>. URL <https://hess.copernicus.org/articles/24/1081/2020/>.
 - [51] Kun Zhang, Mitul Luhar, Manuela I. Brunner, and Anthony J. Parolari. Stream-flow prediction in poorly gauged watersheds in the united states through data-driven sparse sensing. *Water Resources Research*, 59(4):e2022WR034092, 2023. doi: <https://doi.org/10.1029/2022WR034092>. URL <https://agupubs.onlinelibrary.wiley.com/doi/abs/10.1029/2022WR034092>.
 - [52] Günter Blöschl, Julia Hall, Juraj Parajka, Rui AP Perdigão, Bruno Merz, Berit Arheimer, Giuseppe T Aronica, Ardian Bilibashi, Ognjen Bonacci, Marco Borga, et al. Changing climate shifts timing of european floods. *Science*, 357(6351): 588–590, 2017. doi: <https://doi.org/10.1126/science.aan2506>. URL <https://doi.org/10.1126/science.aan2506>.
 - [53] Louise Slater, Georgios Blougouras, Liangkun Deng, Qimin Deng, Emma Ford, Anne Hoek van Dijke, Feini Huang, Shijie Jiang, Yinxue Liu, Simon Moulds, et al. Challenges and opportunities of ml and explainable ai in large-sample hydrology. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 383(2302), 2025. doi: <https://doi.org/10.1098/rsta.2024.0287>. URL <https://doi.org/10.1098/rsta.2024.0287>.
 - [54] Johannes Stephan, Oliver Stegle, and Andreas Beyer. A random forest approach to capture genetic effects in the presence of population structure. *Nature Communications*, 6(1):7432, 2015. doi: <https://doi.org/10.1038/ncomms8432>. URL <https://doi.org/10.1038/ncomms8432>.
 - [55] Yalan Song, Kamlesh Sawadekar, Jonathan M Frame, Ming Pan, Martyn P Clark, Wouter JM Knoben, Andrew W Wood, Kathryn E Lawson, Trupesh Patel, and Chaopeng Shen. Physics-informed, differentiable hydrologic models for capturing unseen extreme events. *Water Resources Research*, 62(2):e2025WR040414, 2026. doi: <https://doi.org/10.1029/2023WR035704>. URL <https://doi.org/10.1029/2023WR035704>.
 - [56] Wenzhong Li, Chengshuai Liu, Caihong Hu, Chaojie Niu, Runxi Li, Ming Li, Yingying Xu, and Lu Tian. Application of a hybrid algorithm of lstm and transformer based on random search optimization for improving rainfall-runoff simulation. *Scientific Reports*, 14(1):11184, 2024. doi: <https://doi.org/10.1038/s41598-024-62127-7>. URL <https://doi.org/10.1038/s41598-024-62127-7>.

- [57] Siyuan Tian, Luigi J. Renzullo, Julien Lerat, and Robert C. Pipunic. Spatio-temporal model variance-covariance approach to assimilating streamflow observations into a distributed landscape water balance model. *Water Resources Research*, 58(4):e2021WR031649, 2022. doi: <https://doi.org/10.1029/2021WR031649>. URL <https://agupubs.onlinelibrary.wiley.com/doi/abs/10.1029/2021WR031649>.
- [58] Khushi Kalasampath, K. N. Spoorthi, Sreeparvathy Sajeev, Sahil Sarma Kuppaa, Kavya Ajay, and Angulakshmi Maruthamuthu. A literature review on applications of explainable artificial intelligence (xai). *IEEE Access*, 13:41111–41140, 2025. doi: <https://doi.org/10.1109/ACCESS.2025.3546681>. URL <https://doi.org/10.1109/ACCESS.2025.3546681>.
- [59] Yoshiyasu Takefuji. Reevaluating feature importance in machine learning: concerns regarding shap interpretations in the context of the eu artificial intelligence act. *Water Research*, 280:123514, 2025. doi: <https://doi.org/10.1016/j.watres.2025.123514>. URL <https://www.sciencedirect.com/science/article/pii/S0043135425004270>.
- [60] L. T. Pham, L. Luo, and A. Finley. Evaluation of random forests for short-term daily streamflow forecasting in rainfall- and snowmelt-driven watersheds. *Hydrology and Earth System Sciences*, 25(6):2997–3015, 2021. doi: <https://doi.org/10.5194/hess-25-2997-2021>. URL <https://hess.copernicus.org/articles/25/2997/2021/>.
- [61] Mahdi Abbasi, Ashkan Farokhnia, Masoud Bahreinimotlagh, and Reza Roozbahani. A hybrid of random forest and deep auto-encoder with support vector regression methods for accuracy improvement and uncertainty reduction of long-term streamflow prediction. *Journal of Hydrology*, 597:125717, 2021. doi: <https://doi.org/10.1016/j.jhydrol.2020.125717>. URL <https://www.sciencedirect.com/science/article/pii/S0022169420311781>.
- [62] Zachary C Lipton. The mythos of model interpretability: In machine learning, the concept of interpretability is both important and slippery. *Queue*, 16(3): 31–57, 2018. doi: <https://doi.org/10.1145/3236386.3241340>. URL <https://doi.org/10.1145/3236386.3241340>.
- [63] Judea Pearl. Theoretical impediments to machine learning with seven sparks from the causal revolution, 2018. URL <https://arxiv.org/abs/1801.04016>.
- [64] Mateusz Pach, Shyamgopal Karthik, Quentin Bouniot, Serge Belongie, and Zeynep Akata. Sparse autoencoders learn monosemantic features in vision-language models, 2025. URL <https://arxiv.org/abs/2504.02821>.
- [65] J. J. Hamman, B. Nijssen, T. J. Bohn, D. R. Gergel, and Y. Mao. The variable infiltration capacity model version 5 (vic-5): infrastructure improvements for new applications and reproducibility. *Geoscientific Model Development*, 11(8):

- 3481–3496, 2018. doi: <https://doi.org/10.5194/gmd-11-3481-2018>. URL <https://gmd.copernicus.org/articles/11/3481/2018/>.
- [66] I. Elizabeth Kumar, Suresh Venkatasubramanian, Carlos Scheidegger, and Sorelle Friedler. Problems with shapley-value-based explanations as feature importance measures, 2020. URL <https://arxiv.org/abs/2002.11097>.
- [67] Riccardo Guidotti, Anna Monreale, Salvatore Ruggieri, Franco Turini, Fosca Giannotti, and Dino Pedreschi. A survey of methods for explaining black box models. *ACM Comput. Surv.*, 51(5), 2018. doi: <https://doi.org/10.1145/3236009>. URL <https://doi.org/10.1145/3236009>.
- [68] Jiangtao Liu, Yuchen Bian, Kathryn Lawson, and Chaopeng Shen. Probing the limit of hydrologic predictability with the transformer network. *Journal of Hydrology*, 637:131389, 2024. doi: <https://doi.org/10.1016/j.jhydrol.2024.131389>. URL <https://www.sciencedirect.com/science/article/pii/S0022169424007844>.
- [69] Louis Dumontet, So-Ra Han, Jun Hyuck Lee, Tae-Jin Oh, and Mingon Kang. Trustworthy prediction of enzyme commission numbers using a hierarchical interpretable transformer. *Nature Communications*, 17(1):1146, 2026. doi: <https://doi.org/10.1038/s41467-026-68727-3>. URL <https://doi.org/10.1038/s41467-026-68727-3>.
- [70] Christos Louizos, Max Welling, and Diederik P. Kingma. Learning sparse neural networks through l_0 regularization, 2018. URL <https://arxiv.org/abs/1712.01312>.
- [71] Bryan Lim, Sercan Ö. Arık, Nicolas Loeff, and Tomas Pfister. Temporal fusion transformers for interpretable multi-horizon time series forecasting. *International Journal of Forecasting*, 37(4):1748–1764, 2021. doi: <https://doi.org/10.1016/j.ijforecast.2021.03.012>. URL <https://www.sciencedirect.com/science/article/pii/S0169207021000637>.
- [72] Michael Zhu and Suyog Gupta. To prune, or not to prune: exploring the efficacy of pruning for model compression, 2017. URL <https://arxiv.org/abs/1710.01878>.
- [73] Hattie Zhou, Janice Lan, Rosanne Liu, and Jason Yosinski. Deconstructing lottery tickets: Zeros, signs, and the supermask. In *Advances in Neural Information Processing Systems*, volume 32, 2019. doi: <https://dl.acm.org/doi/10.5555/3454287.3454610>. URL https://proceedings.neurips.cc/paper_files/paper/2019/file/1113d7a76ffcecalbb350bfe145467c6-Paper.pdf.
- [74] M.A. Fedora and R.L. Beschta. Storm runoff simulation using an antecedent precipitation index (api) model. *Journal of Hydrology*, 112(1):121–133, 1989. doi: [https://doi.org/10.1016/0022-1694\(89\)90184-4](https://doi.org/10.1016/0022-1694(89)90184-4). URL <https://www.sciencedirect.com/science/article/pii/S0022169489901844>.

- [75] Jasper Snoek, Hugo Larochelle, and Ryan P Adams. Practical bayesian optimization of machine learning algorithms. In *Advances in Neural Information Processing Systems*, volume 25, 2012. doi: <https://dl.acm.org/doi/10.5555/2999325.2999464>. URL https://proceedings.neurips.cc/paper_files/paper/2012/file/05311655a15b75fab86956663e1819cd-Paper.pdf.