

Beyond Feature Attribution: Are Deep Learning Models Learning Hydrological Processes?

Supplementary Information

Contents

1	Robust Predictive Performance Across Hydrological Stations	1
2	Specific Methodological and Data Details	2
2.1	Sparse Transformer	2
2.1.1	Transformer Formulation	2
2.1.2	Structured Gating Mechanism	2
2.1.3	Sparse Training Strategy	2
2.2	Interpretability Analysis	3
2.3	Evaluation Metrics	5
2.4	Data	5
3	Experimental Design and Detailed Settings	7
3.1	Learned Core Circuit Structures Reflect Hydrological Knowledge Across Basins	7
3.2	Accurate Volume Response but Limited Sensitivity to Temporal and State Perturbations	8
3.3	STEH-Derived Interpretability Enhances Random Forest Predictive Performance	10
4	Model Hyperparameters	11
4.1	STEH	11
4.2	Random Forest	11

1 Robust Predictive Performance Across Hydrological Stations

Interpretability analysis is most meaningful when the underlying predictions are reasonably accurate. We therefore first verify that STEH achieves reliable predictive performance across hydrological stations[1]. Predictive skill is evaluated using the NSE, and we follow the commonly used threshold of $NSE > 0.5$ to indicate satisfactory performance[2].

As shown in Fig. S1, STEH achieves satisfactory skill for a large proportion of stations, even though it is designed to prioritize interpretability rather than purely optimizing predictive accuracy. Specifically, 327 out of 521 stations (approximately 62.8%) achieve NSE values greater than 0.5[3].




Fig. S1: **Model performance across 521 streamflow gauging stations.** (a) Spatial distribution of NSE, with warmer colors indicating higher predictive skill. (b) Distribution of NSE values across all stations, showing that the majority achieve satisfactory performance ($NSE > 0.5$). (c) Taylor diagram summarizing model performance in terms of correlation, normalized standard deviation, and centered root-mean-square error relative to observations.

These results indicate that the model captures key hydrological dynamics in many basins[4], pro-

viding a practical foundation[1] for the subsequent interpretability analyses. The insights reported below are therefore grounded in meaningful predictive behavior rather than artifacts of uniformly poor performance[5]. We next investigate the learned circuit structures and pathways.

2 Specific Methodological and Data Details

2.1 Sparse Transformer

2.1.1 Transformer Formulation

The Transformer architecture has recently been introduced into hydrological time-series modelling because its self-attention mechanism can effectively capture long-range temporal dependencies in hydro-meteorological sequences[6]. In streamflow prediction, antecedent rainfall, temperature, and soil moisture conditions jointly influence discharge responses across multiple time lags, making long-term dependency modelling particularly important[7].

Let $\mathbf{X} \in \mathbb{R}^{L \times F}$ denote the input sequence with length L and feature dimension F . The self-attention operation is defined as[8]

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax} \left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_k}} \right) \mathbf{V}, \quad (\text{S1})$$

where $\mathbf{Q}$, $\mathbf{K}$, and $\mathbf{V}$ represent the query, key, and value matrices derived from the input representations, and d_k denotes the key dimension used for scaling.

Hidden representations are updated through stacked Transformer encoder blocks[9]:

$$\mathbf{H}^{(l+1)} = \text{Block}^{(l)}(\mathbf{H}^{(l)}), \quad l = 0, \dots, N - 1, \quad (\text{S2})$$

where $\mathbf{H}^{(l)}$ represents the hidden state of layer l , and N denotes the number of encoder layers.

2.1.2 Structured Gating Mechanism

To enable circuit-level interpretability, structured gating mechanisms are introduced at three levels of the network architecture: input variables, attention heads, and feed-forward neurons[10]. The effect of gate-based circuit selection is illustrated in Fig. S2(e–f).

For the input layer, the gated representation is defined as

$$\tilde{\mathbf{X}} = \mathbf{X} \odot \mathbf{m}_{\text{in}}, \quad (\text{S3})$$

where $\mathbf{m}_{\text{in}} \in \{0, 1\}^F$ is a binary mask indicating whether each input feature is retained.

Similarly, head-level and neuron-level masks are defined for each Transformer layer, denoted by $\mathbf{m}_{\text{head}}^{(l)}$ and $\mathbf{m}_{\text{neu}}^{(l)}$, respectively. These masks control the activation of attention heads and feed-forward neurons[11], enabling structured sparsification of the model.

2.1.3 Sparse Training Strategy

Training is conducted in three consecutive stages to stabilize optimization and avoid premature pruning of potentially important connections[12].

Stage I (Origin Model Training). The model is first trained without sparsity constraints to establish a stable predictive baseline (Fig. S2(a–b)).

Stage II (Weight Sparsification). Top- K sparsity is progressively imposed on weight tensors with an annealed sparsity schedule[13] (Fig. S2(c–d)):

$$\rho_t = \rho_{\max} - (\rho_{\max} - \rho_{\min}) \left(1 - \frac{t}{T} \right)^\gamma, \quad (\text{S4})$$

where $\rho_{\min}$ and $\rho_{\max}$ denote the initial and final sparsity ratios. t represents the pruning iteration, T is the total number of steps over which sparsification is applied, and γ is a shape parameter that controls the annealing schedule. Specifically, larger γ values result in a slower increase in sparsity at early stages and a more rapid transition near the end, while smaller γ leads to a more uniform sparsification process.

For each weight group $\mathbf{W}_g$ with N_g elements, the annealed sparsity ratio ρ_t determines the number of retained weights as

$$k_{g,t} = \max(1, \lfloor (1 - \rho_t)N_g \rfloor). \quad (\text{S5})$$

The corresponding magnitude threshold is the $k_{g,t}$ -th largest absolute value:

$$\theta_{g,t} = \text{TopKAbs}(\mathbf{W}_g, k_{g,t})_{\min}, \quad (\text{S6})$$

Stage III (Mask Learning). Learnable gates are optimized with sparsity regularization (Fig. S2(e–f))[14]:

$$\begin{cases} \mathcal{L} = \mathcal{L}_{\text{task}} + \mathcal{L}_{\text{gate}}, \\ \mathcal{L}_{\text{task}} = \frac{1}{M} \sum_{n=1}^M (\hat{y}_n - y_n)^2, \\ \mathcal{L}_{\text{gate}} = \lambda_{\text{gate}} \frac{1}{|\mathcal{G}|} \sum_{g \in \mathcal{G}} p_g. \end{cases} \quad (\text{S7})$$

where M is the number of training samples in a mini-batch, $\mathcal{L}_{\text{task}}$ is the mean squared error loss, $\mathcal{L}_{\text{gate}}$ is the overall gate sparsity regularizer, $\mathcal{G}$ denotes all learnable gates. We consider two types of gates: input-level gates and node-level gates, corresponding to *input* and *head_neu*, respectively, and p_g is the activation probability of gate g .

After training, a compact core predictive circuit is extracted by thresholding gate probabilities[15]. Let $p_i = \sigma(z_i)$ denote the activation probability of gate i . A unit is retained if

$$a_i = \mathbb{I}(p_i > \tau), \quad (\text{S8})$$

where τ is the pruning threshold, which is optimized by BO algorithm. Applying this rule to all gates yields a compact model containing only the most important computational pathways.

2.2 Interpretability Analysis

(1) Ablation-based Importance

To quantify the structural importance of individual computational units, targeted ablation is applied to each candidate unit u [16]. The importance score is defined as

$$\Delta \text{MSE}_u = \text{MSE}(\text{ablate}(u)) - \text{MSE}(\text{base}). \quad (\text{S9})$$

Unlike gradient-based attribution methods, this intervention-based measurement directly evaluates the structural contribution of each component by observing the performance change after structural removal[17].

(2) Interaction Analysis

To analyze dependencies between components, pairwise interaction scores are defined as

$$I_{ij} = \Delta_{ij} - \Delta_i - \Delta_j. \quad (\text{S10})$$

Positive values indicate synergistic interactions, whereas negative values suggest redundancy or functional substitution.




Fig. S2: Three-stage sparsification framework for predictive circuit discovery in Transformer. The model is trained through three stages to progressively transform a dense Transformer into a sparse predictive circuit. (a–b) Stage I: Origin model training, where the Transformer is trained without sparsity constraints and all input features contribute to the prediction. (c–d) Stage II: Weight sparsification progressively prunes redundant connections according to the sparsity schedule. (e–f) Stage III: Mask learning optimizes learnable gates so that only a subset of input features contributes to the final prediction. For clarity, only a subset of connections between layers is illustrated in the figure, while the actual model remains fully connected, and decoder attention layers are omitted for simplicity. Note that the proposed method does not explicitly modify the network architecture; sparsity is achieved by setting weights or gates to zero, while the visual removal of units in the figure is only for illustrative purposes. The number of input variables shown in the figure is fewer than the actual input feature set; this simplification is adopted solely for visual clarity.

(3) Top-K Faithfulness

To evaluate the trade-off between circuit compactness and predictive fidelity, a top- K faithfulness analysis is conducted by retaining only the most important units and measuring the resulting performance degradation[18].

(4) Physical Probing

To examine whether hidden representations encode physically meaningful hydrological information, linear probes are applied to hidden states[19].

Let $\mathbf{h}_n \in \mathbb{R}^D$ denote the hidden representation of sample n . The probe model is defined as

$$\hat{y}_n^{\text{phys}} = \mathbf{w}^\top \mathbf{h}_n + b, \quad (\text{S11})$$

where $\mathbf{w}$ and b are parameters estimated using least squares regression.

Probe performance is evaluated using the coefficient of determination (R^2). To avoid probe overfitting, the probe model is intentionally restricted to a simple linear form without hidden layers.

Representative hydrological variables such as cumulative rainfall[20], antecedent precipitation index (API)[21], and rainfall change rate are used as probe targets. Let P_t denote precipitation at day t , and let

W be the look-back window length. Their definitions are

$$R_t^{\text{cum}}(W) = \sum_{k=0}^{W-1} P_{t-k}, \quad (\text{S12})$$

$$\text{API}_t(W, \kappa) = \sum_{k=0}^{W-1} \kappa^k P_{t-k}, \quad 0 < \kappa < 1, \quad (\text{S13})$$

$$R_t^{\text{chg}} = \frac{P_t - P_{t-1}}{P_{t-1} + \varepsilon}, \quad (\text{S14})$$

where κ is a decay factor controlling the memory of antecedent rainfall and ε is a small constant to avoid division by zero.

2.3 Evaluation Metrics

Model evaluation is conducted from three perspectives: predictive accuracy, explainability properties, and hydro-climatic diagnostics.

Predictive accuracy is measured using standard hydrological metrics including RMSE[3], MAE[22], and NSE[4]. Explainability-related metrics include ablation importance (ΔMSE , Eq. S9), interaction score (I_{ij} , Eq. S10), and top- K faithfulness[18] trends. Structural sparsity statistics are also reported to quantify circuit compactness.

Hydro-climatic diagnostics are used to characterize the climatic regimes of catchments. Let $\bar{P}$, $\overline{\text{PET}}$, and $\bar{Q}$ denote the long-term mean annual precipitation, potential evapotranspiration, and streamflow, respectively, and let P_m ($m = 1, \dots, 12$) be the mean monthly precipitation with $\bar{P}_{\text{ann}} = \sum_{m=1}^{12} P_m$.

The **aridity index** (AI)[23] describes the balance between water supply and atmospheric demand:

$$\text{AI} = \frac{\bar{P}}{\overline{\text{PET}}}. \quad (\text{S15})$$

Values $\text{AI} < 1$ indicate arid conditions, while $\text{AI} > 1$ indicates humid conditions.

The **Streamflow coefficient** (C_r)[24] represents the fraction of precipitation that becomes streamflow:

$$C_r = \frac{\bar{Q}}{\bar{P}}. \quad (\text{S16})$$

The **precipitation seasonality index** (SI)[25] quantifies the degree of seasonal concentration of rainfall:

$$\text{SI} = \frac{1}{\bar{P}_{\text{ann}}} \sum_{m=1}^{12} \left| P_m - \frac{\bar{P}_{\text{ann}}}{12} \right|, \quad (\text{S17})$$

where $\text{SI} = 0$ indicates perfectly uniform monthly distribution and larger values correspond to stronger seasonal concentration. Together, these three diagnostics are used to stratify stations by hydro-climatic regime and to contextualize model performance across different environmental conditions.

2.4 Data

Because streamflow record availability varies across stations, we apply a unified quality-control procedure within this period. Stations with more than 100 missing daily values in the target streamflow series are excluded, whereas stations with at most 100 missing days are retained and missing values are filled by linear interpolation. After preprocessing, 521 stations remain for subsequent analysis, as Fig S3 shows. All retained stations provide multi-year continuous records over the common window (well beyond 20 years), ensuring sufficient sample size for deep learning training.

The model input variables are summarized in Table S1. For each station, seven daily meteorological forcing variables from Daymet are used, and daily streamflow is additionally incorporated to represent antecedent hydrological state knowledge.




Fig. S3: Spatial distribution of selected streamflow stations after quality control

Table S1: Selected Input Meteorological and Hydrological Variables

Group	Feature name	Unit
Energy	day length (dayl)	s day ⁻¹
	shortwave radiation (srad)	W m ⁻²
Temperature	maximum air temperature (tmax)	°C
	minimum air temperature (tmin)	°C
Atmospheric moisture	water vapour pressure (vp)	Pa
Precipitation	precipitation (prcp)	mm day ⁻¹
Snow	snow water equivalent (swe)	mm
Runoff	streamflow (Q)	mm day ⁻¹

3 Experimental Design and Detailed Settings

3.1 Learned Core Circuit Structures Reflect Hydrological Knowledge Across Basins

To analyze how learned interaction structures in core circuits vary under different hydro-climatic conditions, we constructed multiple climate classification indicators based on static basin attributes, including aridity class, snow regime, and geo-climatic type[5]. These classifications were used for post-hoc grouping and statistical analysis of interaction structures across basins.

Aridity Class

Dryness conditions were classified using the aridity index (`aridity`) through binning[23]. The classification criteria are summarized in Table S2.

Table S2: Aridity-based hydro-climatic classification

Aridity range	Class	Description
$-1 \leq aridity < 0.65$	Humid	Humid
$0.65 \leq aridity < 1.0$	Sub-humid	Sub-humid
$aridity \geq 1.0$	Arid	Arid

The thresholds (0.65 and 1.0) follow established hydro-climatic stratification practice and effectively distinguish energy-limited from water-limited basins[26]. They provide physically meaningful groupings for subsequent analysis of interaction structure distributions.

Snow Regime

Snow regime was classified based on snowfall fraction (`frac_snow`) to distinguish rain-dominated, mixed, and snow-influenced basins, as summarized in Table S3[27].

Table S3: Snow regime classification based on snowfall fraction

Fraction range	Class	Description
$-0.1 \leq frac_snow < 0.1$	Rain-dominant	Rain-dominated
$0.1 \leq frac_snow < 0.3$	Mixed	Mixed
$frac_snow \geq 0.3$	Snow-influenced	Snow-influenced

The selected thresholds (0.1 and 0.3) reflect the gradual transition from rain-driven to snow-driven runoff. Low snowfall fractions correspond to basins dominated by direct rainfall, whereas higher fractions indicate increasing importance of snow accumulation and melt processes[28].

Geo-climatic Type

Geo-climatic type was determined using a hierarchical rule based on elevation (`forcing_elev_m`), longitude (`lon`), and aridity[26]. The pseudocode is shown in Algorithm 1.

This hierarchical classification first identifies high-elevation regions ($\geq 1500m$), then distinguishes oceanic versus inland climates based on longitude, and finally separates arid continental basins using a higher aridity threshold (1.5). This approach provides a physically interpretable stratification of basins across major hydro-climatic regimes.

Algorithm 1 Hierarchical classification for geo-climatic type

Require: Elevation *forcing_elev_m*, Longitude *lon*, Aridity *aridity*

Ensure: Geo-climatic type *geo_climate_type*

```
1: if forcing_elev_m  $\geq$  1500 then
2:   geo_climate_type  $\leftarrow$  "Plateau/Mountain"
3: else if lon  $\leq$  -120 then
4:   geo_climate_type  $\leftarrow$  "Oceanic"
5: else if aridity  $\geq$  1.5 then
6:   geo_climate_type  $\leftarrow$  "Arid Continental"
7: else
8:   geo_climate_type  $\leftarrow$  "Humid Continental/Subtropical"
9: end if
```

3.2 Accurate Volume Response but Limited Sensitivity to Temporal and State Perturbations

To examine whether the model encodes physically consistent hydrological sensitivities—rather than relying on spurious statistical correlations—we designed a station-wise counterfactual sensitivity framework on the test split[29]. Let f_θ denote the trained sparse transformer, $\mathbf{X}_i \in \mathbb{R}^{N_i \times L \times F}$ the input tensor for station i (with N_i samples, lookback length L , and F predictors), and $\mathbf{y}_i \in \mathbb{R}^{N_i}$ the observed discharge target. For each intervention setting s , we define

$$\hat{\mathbf{y}}_i^{(0)} = f_\theta(\mathbf{X}_i), \quad \hat{\mathbf{y}}_i^{(s)} = f_\theta(\mathcal{T}_s(\mathbf{X}_i)), \quad (\text{S18})$$

where $\mathcal{T}_s(\cdot)$ denotes an intervention operator applied in physical space and subsequently mapped back to the normalized model space using the same normalization statistics as in training[30]. This paired design isolates the model’s response to controlled perturbations while eliminating cross-sample confounding effects[31].

Experiment design

We considered four intervention families to disentangle model sensitivities to hydrological drivers, including volume, temporal distribution, catchment state, and thermal forcing[32]. The perturbation ranges are chosen to represent moderate yet plausible deviations that remain within the support of the training distribution.

(1) Rainfall scaling. Let P_t denote precipitation at lag $t \in \{1, \dots, L\}$. With scaling factor α , the intervention is

$$P'_t = \alpha P_t, \quad \alpha \in \{0.8, 0.9, 1.1, 1.2\}. \quad (\text{S19})$$

This modifies storm magnitude while preserving the original temporal structure of rainfall[33].

(2) Rainfall pattern redistribution. Let $\mathbf{w}^{(m)} = (w_1^{(m)}, \dots, w_L^{(m)})$ be a mode-specific nonnegative weight vector (front-loaded, middle-loaded, back-loaded, uniform) satisfying $\sum_{t=1}^L w_t^{(m)} = 1$. The redistributed rainfall is

$$P'_t = \left(\sum_{j=1}^L P_j \right) w_t^{(m)}. \quad (\text{S20})$$

This conserves total rainfall volume while altering its temporal distribution within the input window[34]. We note that this intervention represents an idealized redistribution rather than a physically simulated storm evolution, and is intended to isolate model sensitivity to intra-window timing patterns.

(3) Streamflow-input scaling (state perturbation). Let Q_t^{in} denote the lagged streamflow input channel. With scaling factor β ,

$$Q_t^{\text{in}'} = \beta Q_t^{\text{in}}, \quad t = 1, \dots, L, \quad \beta \in \{0.1, 0.5, 0.9, 1.1, 1.5, 1.9\}. \quad (\text{S21})$$

This isolates sensitivity to hydrological state magnitude in the full lookback window while preserving the original temporal pattern[35].

(4) Temperature increase perturbation. For each temperature-related predictor T_t , we apply

$$T_t' = T_t + \Delta T, \quad \Delta T \in \{3, 5, 7\}^\circ\text{C}. \quad (\text{S22})$$

This probes thermal forcing effects (e.g., evapotranspiration or snowmelt processes), particularly in energy-limited or snow-influenced catchments[36].

In total, the intervention set contains $4 + 4 + 6 + 3 = 17$ scenarios per station. For each station and scenario, responses are computed against the same baseline prediction, enabling paired comparison[37].

Evaluation metrics

For station i and scenario s , with sample-wise observations $y_{i,n}$ and predictions $\hat{y}_{i,n}^{(s)}$ ($n = 1, \dots, N_i$), we evaluate the following metrics[2, 3]: **(1) MSE.**

(2) Predicted streamflow volume.

$$V_i^{(s)} = \sum_{n=1}^{N_i} \hat{y}_{i,n}^{(s)}. \quad (\text{S23})$$

(3) Predicted peak timing index.

$$\tau_i^{(s)} = \arg \max_{1 \leq n \leq N_i} \hat{y}_{i,n}^{(s)}. \quad (\text{S24})$$

Peak-timing-based diagnostics are commonly used to characterize flow timing shifts under altered hydro-climatic forcing[38].

Relative to baseline $s = 0$, we define

$$\Delta \text{MSE}_i^{(s)} = \text{MSE}_i^{(s)} - \text{MSE}_i^{(0)}, \quad \Delta V_i^{(s)} = V_i^{(s)} - V_i^{(0)}, \quad \Delta \tau_i^{(s)} = \tau_i^{(s)} - \tau_i^{(0)}. \quad (\text{S25})$$

For visualization comparability across stations, relative changes are reported as

$$\delta \text{MSE}_i^{(s)}(\%) = 100 \times \frac{\Delta \text{MSE}_i^{(s)}}{\max(|\text{MSE}_i^{(0)}|, 1)}, \quad \delta V_i^{(s)}(\%) = 100 \times \frac{\Delta V_i^{(s)}}{\max(|V_i^{(0)}|, 1)}. \quad (\text{S26})$$

Physical consistency criteria

To assess whether model responses align with expected hydrological behavior, we adopt direction-based consistency criteria[30]:

For rainfall scaling and streamflow-input scaling, expected behavior requires directionally consistent changes in streamflow volume, i.e., ΔV increases for scaling factors greater than 1 and decreases for factors less than 1[33].

For rainfall pattern redistribution and temperature perturbation, expected behavior requires that model performance does not exhibit systematic degradation, i.e., no consistent increase in prediction error across scenarios[32].

Finally, for each intervention family, station-level expectedness is aggregated by majority vote over its scenarios. Statistical significance of paired scenario effects is assessed by two-sided Wilcoxon signed-rank tests, with Benjamini–Hochberg correction for multiple comparisons.

3.3 STEH-Derived Interpretability Enhances Random Forest Predictive Performance

To evaluate whether the hydrological circuit structures revealed by STEH can improve a conventional machine-learning benchmark, we designed a paired station-wise comparison between a *baseline* Random Forest (RF) and a *STEH-informed modified* RF[39]. Both models are trained and tested on the same station, temporal split, and target variable, ensuring that any performance differences can be attributed to the STEH-derived design rather than inconsistencies in data or experimental setup[40].

Let station index be $i \in \{1, \dots, M\}$, predictors $\mathbf{X}_i$, and discharge target $\mathbf{y}_i$. For each station, we train

$$\hat{\mathbf{y}}_i^{\text{base}} = g_{\phi_i}(\mathbf{X}_i), \quad \hat{\mathbf{y}}_i^{\text{mod}} = g_{\psi_i}(\tilde{\mathbf{X}}_i), \quad (\text{S27})$$

where g_{ϕ_i} and g_{ψ_i} denote RF predictors with station-specific fitted parameters.

Experiment design

During model training, the hyperparameters of both Random Forest models are optimized using Bayesian optimisation on the validation set[41], ensuring that performance differences are not attributable to discrepancies in tuning strategy. Specifically, the tuned hyperparameters include tree depth ($\in [2, 8]$), the minimum number of samples required for node splitting ($\in [2, 8]$), the minimum number of samples at leaf nodes ($\in [1, 4]$), and the feature sampling ratio for split selection ($\in [0.5, 0.8]$), further information can be seen in Supplementary Notes 4.2.

The task is formulated as a one-step-ahead prediction problem, where a fixed 10-day input window is used to predict next-day discharge. Both baseline and modified models use the same input length ($L = 10$). Data are split chronologically into training, validation, and testing sets with ratios of 0.7/0.15/0.15, respectively, on a per-station basis. Both models use identical raw drivers, including precipitation, temperature, shortwave radiation, and antecedent discharge.

Baseline RF. The baseline model adopts standard lagged inputs and rolling statistical summaries, without incorporating explicit hydrological regime information.

STEH-informed modified RF. The modified model incorporates structural insights revealed by STEH (primarily P-Q, Q-Q, and T-Q pathways). Specifically, it includes physically interpretable aggregated features (e.g., $P_3, P_7, P_{14}, Q_3, Q_7, Q_3 - Q_7, T_7^+, T_7^-$), introduces regime indicators to distinguish dominant hydrological processes (P-Q, Q-Q, T-Q), and applies light regime-aware sample weighting for selected regimes (P-Q and T-Q) during training.

The dominant regime for each station is identified over the training period based on correlation analysis[34]. Specifically, we compute

$$r_{\text{PQ}} = \text{corr}(P_t, Q_t), \quad r_{\text{QQ}} = \text{corr}(Q_{t-1}, Q_t), \quad r_{\text{TQ}} = \text{corr}(T_t, Q_t), \quad (\text{S28})$$

and define the regime as

$$\mathcal{R}_i = \arg \max \{r_{\text{PQ}}, r_{\text{QQ}}, r_{\text{TQ}}\}. \quad (\text{S29})$$

This regime label is then used as part of the model input.

Training and evaluation protocol

Model performance is evaluated using Nash-Sutcliffe efficiency (NSE), mean squared error (MSE), and Kling-Gupta efficiency (KGE), with test NSE serving as the primary comparison metric[2]. Evaluation is conducted on stations where the sparse-transformer benchmark achieves NSE greater than 0.5[4].

For the selected stations, we report the mean NSE, mean MSE and mean KGE of both baseline and modified RF models, and perform a paired t-test on station-wise NSE differences.

4 Model Hyperparameters

4.1 STEH

Hyperparameters were systematically optimized within the predefined search spaces listed in Table S4[41], ensuring reproducibility and minimizing subjective bias. The sparsity-related settings are consistent with prior sparse-Transformer practice[15].

Due to the large number of study sites, reporting all optimized hyperparameter configurations in this Supplementary Information is impractical. Therefore, the complete set of best hyperparameters is provided in an repository for full transparency and reproducibility: https://github.com/Polyethylenekjc/Sparse-Transformer-In-Streamflow-Prediction/blob/main/outcomes/STEH_Best_Hyperparams.csv

Table S4: Hyperparameter settings and search space

Hyperparameter	Description	Search Space
Learning rate	Controls optimisation step size	$10^{-4} - 10^{-2}$
Batch size	Number of samples per mini-batch	32 – 256
Hidden size	Dimension of hidden representations	16 – 256
Number of layers	Number of Transformer encoder layers	1 – 16
Weight sparsity	Target fraction of pruned weights in Stage II	0.7 – 0.95
Activation sparsity	Sparsification applied on the last activation dimension.	0.6 – 0.9
Gate regularisation	Strength of gate sparsity regularization	$10^{-5} - 10^{-3}$

4.2 Random Forest

Hyperparameters of the Random Forest models were systematically optimized using Bayesian optimisation within the predefined search spaces listed in Table S5[41], ensuring reproducibility and minimizing subjective bias.

Due to the large number of study sites, reporting all optimized hyperparameter configurations in this Supplementary Information is impractical. Therefore, the complete set of best hyperparameters is provided in a repository for full transparency and reproducibility: https://github.com/Polyethylenekjc/Sparse-Transformer-In-Streamflow-Prediction/blob/main/outcomes/Rf_Best_Hyperparams.csv

Table S5: Hyperparameter settings and search space for Random Forest

Hyperparameter	Description	Search Space
Tree depth	Maximum tree depth	2 – 8
Min samples (split)	Minimum number of samples required to split	2 – 8
Min samples (leaf)	Minimum number of samples per leaf	1 – 4
Feature ratio	Fraction of features considered per split	0.5 – 0.8

Note: “Baseline” refers to the standard Random Forest model without incorporating STEH-derived hydrological knowledge, while “Modified” denotes the Random Forest model augmented with STEH-informed structured features.

References

- [1] Riccardo Guidotti, Anna Monreale, Salvatore Ruggieri, Franco Turini, Fosca Giannotti, and Dino Pedreschi. A survey of methods for explaining black box models. *ACM Comput. Surv.*, 51(5), 2018. doi: <https://doi.org/10.1145/3236009>. URL <https://doi.org/10.1145/3236009>.

- [2] Daniel N Moriasi, Jeffrey G Arnold, Michael W Van Liew, Ronald L Bingner, Daren Harmel, and Thomas L Veith. Model evaluation guidelines for systematic quantification of accuracy in watershed simulations. *Transactions of the ASABE*, 50(3):885–900, 2007. doi: <https://doi.org/10.13031/2013.23153>. URL <https://doi.org/10.13031/2013.23153>.
- [3] Hiren Solanki, Urmin Vegad, Anuj Kushwaha, and Vimal Mishra. Improving streamflow prediction using multiple hydrological models and machine learning methods. *Water Resources Research*, 61(1):e2024WR038192, 2025. doi: <https://doi.org/10.1029/2024WR038192>. URL <https://agupubs.onlinelibrary.wiley.com/doi/abs/10.1029/2024WR038192>.
- [4] Kun Zhang, Mitul Luhar, Manuela I. Brunner, and Anthony J. Parolari. Streamflow prediction in poorly gauged watersheds in the united states through data-driven sparse sensing. *Water Resources Research*, 59(4):e2022WR034092, 2023. doi: <https://doi.org/10.1029/2022WR034092>. URL <https://agupubs.onlinelibrary.wiley.com/doi/abs/10.1029/2022WR034092>.
- [5] F. Kratzert, D. Klotz, G. Shalev, G. Klambauer, S. Hochreiter, and G. Nearing. Towards learning universal, regional, and local hydrological behaviors via machine learning applied to large-sample datasets. *Hydrology and Earth System Sciences*, 23(12):5089–5110, 2019. doi: <https://doi.org/10.5194/hess-23-5089-2019>. URL <https://hess.copernicus.org/articles/23/5089/2019/>.
- [6] Jiangtao Liu, Yuchen Bian, Kathryn Lawson, and Chaopeng Shen. Probing the limit of hydrologic predictability with the transformer network. *Journal of Hydrology*, 637:131389, 2024. doi: <https://doi.org/10.1016/j.jhydrol.2024.131389>. URL <https://www.sciencedirect.com/science/article/pii/S0022169424007844>.
- [7] Wenzhong Li, Chengshuai Liu, Caihong Hu, Chaojie Niu, Runxi Li, Ming Li, Yingying Xu, and Lu Tian. Application of a hybrid algorithm of lstm and transformer based on random search optimization for improving rainfall-runoff simulation. *Scientific Reports*, 14(1):11184, 2024. doi: <https://doi.org/10.1038/s41598-024-62127-7>. URL <https://doi.org/10.1038/s41598-024-62127-7>.
- [8] Louis Dumontet, So-Ra Han, Jun Hyuck Lee, Tae-Jin Oh, and Mingon Kang. Trustworthy prediction of enzyme commission numbers using a hierarchical interpretable transformer. *Nature Communications*, 17(1):1146, 2026. doi: <https://doi.org/10.1038/s41467-026-68727-3>. URL <https://doi.org/10.1038/s41467-026-68727-3>.
- [9] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30, 2017. doi: <https://dl.acm.org/doi/10.5555/3295222.3295349>. URL https://proceedings.neurips.cc/paper_files/paper/2017/file/3f5ee243547dee91fbd053c1c4a845aa-Paper.pdf.
- [10] Christos Louizos, Max Welling, and Diederik P. Kingma. Learning sparse neural networks through l_0 regularization, 2018. URL <https://arxiv.org/abs/1712.01312>.
- [11] Bryan Lim, Sercan Ö. Arık, Nicolas Loeff, and Tomas Pfister. Temporal fusion transformers for interpretable multi-horizon time series forecasting. *International Journal of Forecasting*, 37(4):1748–1764, 2021. doi: <https://doi.org/10.1016/j.ijforecast.2021.03.012>. URL <https://www.sciencedirect.com/science/article/pii/S0169207021000637>.
- [12] Song Han, Jeff Pool, John Tran, and William Dally. Learning both weights and connections for efficient neural network. In *Advances in Neural Information Processing Systems*, volume 28, 2015. doi: <https://dl.acm.org/doi/10.5555/2969239>.

2969366. URL https://proceedings.neurips.cc/paper_files/paper/2015/file/ae0eb3eed39d2bcef4622b2499a05fe6-Paper.pdf.
- [13] Michael Zhu and Suyog Gupta. To prune, or not to prune: exploring the efficacy of pruning for model compression, 2017. URL <https://arxiv.org/abs/1710.01878>.
- [14] Hattie Zhou, Janice Lan, Rosanne Liu, and Jason Yosinski. Deconstructing lottery tickets: Zeros, signs, and the supermask. In *Advances in Neural Information Processing Systems*, volume 32, 2019. doi: <https://dl.acm.org/doi/10.5555/3454287.3454610>. URL https://proceedings.neurips.cc/paper_files/paper/2019/file/1113d7a76ffceca1bb350bfe145467c6-Paper.pdf.
- [15] Leo Gao, Achyuta Rajaram, Jacob Coxon, Soham V. Govande, Bowen Baker, and Dan Mossing. Weight-sparse transformers have interpretable circuits, 2025. URL <https://arxiv.org/abs/2511.13653>.
- [16] Ari S. Morcos, David G. T. Barrett, Neil C. Rabinowitz, and Matthew Botvinick. On the importance of single directions for generalization, 2018. URL <https://arxiv.org/abs/1803.06959>.
- [17] Kevin Wang, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt. Interpretability in the wild: a circuit for indirect object identification in gpt-2 small, 2022. URL <https://arxiv.org/abs/2211.00593>.
- [18] Wojciech Samek, Alexander Binder, Grégoire Montavon, Sebastian Lapuschkin, and Klaus-Robert Müller. Evaluating the visualization of what a deep neural network has learned. *IEEE Transactions on Neural Networks and Learning Systems*, 28(11):2660–2673, 2017. doi: <https://doi.org/10.1109/TNNLS.2016.2599820>. URL <https://doi.org/10.1109/TNNLS.2016.2599820>.
- [19] Vikas Verma, Alex Lamb, Christopher Beckham, Amir Najafi, Ioannis Mitliagkas, David Lopez-Paz, and Yoshua Bengio. Manifold mixup: Better representations by interpolating hidden states. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97, pages 6438–6447, 2019. doi: <https://doi.org/10.48550/arXiv.1806.05236>. URL <https://proceedings.mlr.press/v97/verma19a.html>.
- [20] Dina Pirone, Luigi Cimorelli, Giuseppe Del Giudice, and Domenico Pianese. Short-term rainfall forecasting using cumulative precipitation fields from station data: a probabilistic machine learning approach. *Journal of Hydrology*, 617:128949, 2023. doi: <https://doi.org/10.1016/j.jhydrol.2022.128949>. URL <https://www.sciencedirect.com/science/article/pii/S0022169422015190>.
- [21] M.A. Fedora and R.L. Beschta. Storm runoff simulation using an antecedent precipitation index (api) model. *Journal of Hydrology*, 112(1):121–133, 1989. doi: [https://doi.org/10.1016/0022-1694\(89\)90184-4](https://doi.org/10.1016/0022-1694(89)90184-4). URL <https://www.sciencedirect.com/science/article/pii/0022169489901844>.
- [22] Siyuan Tian, Luigi J. Renzullo, Julien Lerat, and Robert C. Pipunic. Spatio-temporal model variance-covariance approach to assimilating streamflow observations into a distributed landscape water balance model. *Water Resources Research*, 58(4):e2021WR031649, 2022. doi: <https://doi.org/10.1029/2021WR031649>. URL <https://agupubs.onlinelibrary.wiley.com/doi/abs/10.1029/2021WR031649>.
- [23] Robert J Zomer, Jianchu Xu, and Antonio Trabucco. Version 3 of the global aridity index and potential evapotranspiration database. *Scientific Data*, 9(1):409, 2022. doi: <https://doi.org/10.1038/s41597-022-01493-1>. URL <https://doi.org/10.1038/s41597-022-01493-1>.

- [24] Shujie Cheng, Lei Cheng, Pan Liu, Shujing Qin, Lu Zhang, Chong-Yu Xu, Lihua Xiong, Liu Liu, and Jun Xia. An analytical baseflow coefficient curve for depicting the spatial variability of mean annual catchment baseflow. *Water Resources Research*, 57(8):e2020WR029529, 2021. doi: <https://doi.org/10.1029/2020WR029529>. URL <https://agupubs.onlinelibrary.wiley.com/doi/abs/10.1029/2020WR029529>.
- [25] R. P. D. Walsh and D. M. Lawler. Rainfall seasonality: Description, spatial patterns and change through time. *Weather*, 36(7):201–208, 1981. doi: <https://doi.org/10.1002/j.1477-8696.1981.tb05400.x>. URL <https://rmets.onlinelibrary.wiley.com/doi/abs/10.1002/j.1477-8696.1981.tb05400.x>.
- [26] F. U. Jehn, K. Bestian, L. Breuer, P. Kraft, and T. Houska. Using hydrological and climatic catchment clusters to explore drivers of catchment behavior. *Hydrology and Earth System Sciences*, 24(3):1081–1100, 2020. doi: <https://doi.org/10.5194/hess-24-1081-2020>. URL <https://hess.copernicus.org/articles/24/1081/2020/>.
- [27] W. R. Berghuijs, R. A. Woods, and M. Hrachowitz. A precipitation shift from snow towards rain leads to a decrease in streamflow. *Nature Climate Change*, 4(7):583–586, 2014. doi: <https://doi.org/10.1038/nclimate2246>. URL <https://doi.org/10.1038/nclimate2246>.
- [28] H. Biemans, C. Siderius, A. F. Lutz, S. Nepal, B. Ahmad, T. Hassan, W. von Bloh, R. R. Wijngaard, P. Wester, A. B. Shrestha, and W. W. Immerzeel. Importance of snow and glacier meltwater for agriculture on the indo-gangetic plain. *Nature Sustainability*, 2(7):594–601, 2019. doi: <https://doi.org/10.1038/s41893-019-0305-3>. URL <https://doi.org/10.1038/s41893-019-0305-3>.
- [29] Filip Simic et al. A comprehensive analysis of perturbation methods in explainable ai. *Scientific Reports*, 15(1):9538, 2025. doi: [10.1038/s41598-025-09538-2](https://doi.org/10.1038/s41598-025-09538-2).
- [30] Jordan S Read, Xiaowei Jia, Jared Willard, Alison P Appling, Jacob A Zwart, Samantha K Oliver, Anuj Karpatne, Gretchen JA Hansen, Paul C Hanson, William Watkins, et al. Process-guided deep learning predictions of lake water temperature. *Water Resources Research*, 55(11):9173–9190, 2019. doi: <https://doi.org/10.1029/2019WR024922>. URL <https://doi.org/10.1029/2019WR024922>.
- [31] Manjula Devak and C. T. Dhanya. Sensitivity analysis of hydrological models: review and way forward. *Journal of Water and Climate Change*, 8(4):557–575, 06 2017. ISSN 2040-2244. doi: [10.2166/wcc.2017.149](https://doi.org/10.2166/wcc.2017.149). URL <https://doi.org/10.2166/wcc.2017.149>.
- [32] M. Jeung, Y. Her, S.-S. Baek, and K. Yoon. Sensitivity of hydrological machine learning prediction accuracy to information quantity and quality. *Hydrology and Earth System Sciences*, 30(4):1077–1095, 2026. doi: [10.5194/hess-30-1077-2026](https://doi.org/10.5194/hess-30-1077-2026). URL <https://hess.copernicus.org/articles/30/1077/2026/>.
- [33] Grey S Nearing, Frederik Kratzert, Alden Keefe Sampson, Craig S Pelissier, Daniel Klotz, Jonathan M Frame, Cristina Prieto, and Hoshin V Gupta. What role does hydrological science play in the age of machine learning? *Water Resources Research*, 57(3):e2020WR028091, 2021. doi: <https://doi.org/10.1029/2020WR028091>. URL <https://doi.org/10.1029/2020WR028091>.
- [34] Samuel E. Franzen, Mozhgan A. Farahani, and Allison E. Goodwell. Information flows: Characterizing precipitation-streamflow dependencies in the colorado headwaters with an information theory approach. *Water Resources Research*, 56(10):e2019WR026133, 2020. doi: <https://doi.org/10.1029/2019WR026133>. URL <https://agupubs.onlinelibrary.wiley.com/doi/abs/10.1029/2019WR026133>.

- [35] Sina Apak, Huseyin Cagan Kilinc, Adem Yurtsever, Hilal Haznedar, and Furkan Ozkan. Multi-resolution adaptive channel fusion transformer encoder lstm for accurate streamflow prediction. *Scientific Reports*, 2026. doi: <https://doi.org/10.1038/s41467-026-70523-y>. URL <https://doi.org/10.1038/s41467-026-70523-y>.
- [36] Paul CD Milly and Krista A Dunne. Colorado river flow dwindles as warming-driven loss of reflective snow energizes evaporation. *Science*, 367(6483):1252–1255, 2020. doi: <https://doi.org/10.1126/science.aay9187>. URL <https://doi.org/10.1126/science.aay9187>.
- [37] Louise Slater, Georgios Blougouras, Liangkun Deng, Qimin Deng, Emma Ford, Anne Hoek van Dijke, Feini Huang, Shijie Jiang, Yinxue Liu, Simon Moulds, et al. Challenges and opportunities of ml and explainable ai in large-sample hydrology. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 383(2302), 2025. doi: <https://doi.org/10.1098/rsta.2024.0287>. URL <https://doi.org/10.1098/rsta.2024.0287>.
- [38] Günter Blöschl, Julia Hall, Juraj Parajka, Rui AP Perdigão, Bruno Merz, Berit Arheimer, Giuseppe T Aronica, Ardian Bilibashi, Ognjen Bonacci, Marco Borga, et al. Changing climate shifts timing of european floods. *Science*, 357(6351):588–590, 2017. doi: <https://doi.org/10.1126/science.aan2506>. URL <https://doi.org/10.1126/science.aan2506>.
- [39] L. T. Pham, L. Luo, and A. Finley. Evaluation of random forests for short-term daily streamflow forecasting in rainfall- and snowmelt-driven watersheds. *Hydrology and Earth System Sciences*, 25(6):2997–3015, 2021. doi: <https://doi.org/10.5194/hess-25-2997-2021>. URL <https://hess.copernicus.org/articles/25/2997/2021/>.
- [40] Mahdi Abbasi, Ashkan Farokhnia, Masoud Bahreinimotlagh, and Reza Roozbahani. A hybrid of random forest and deep auto-encoder with support vector regression methods for accuracy improvement and uncertainty reduction of long-term streamflow prediction. *Journal of Hydrology*, 597:125717, 2021. doi: <https://doi.org/10.1016/j.jhydrol.2020.125717>. URL <https://www.sciencedirect.com/science/article/pii/S0022169420311781>.
- [41] Jasper Snoek, Hugo Larochelle, and Ryan P Adams. Practical bayesian optimization of machine learning algorithms. In *Advances in Neural Information Processing Systems*, volume 25, 2012. doi: <https://dl.acm.org/doi/10.5555/2999325.2999464>. URL https://proceedings.neurips.cc/paper_files/paper/2012/file/05311655a15b75fab86956663e1819cd-Paper.pdf.